

E-commerce shopping dataset

August 30, 2024

0.1 E-commerce Shopping website Analytics

Objective: To conduct a thorough exploratory data analysis (EDA) and hypothesis testing on two comprehensive datasets one containing information on customers visiting the shopping site for purchase and another that has demographic, purchase, and marketing information about the group of people

Expectations:

The project expects you to Analyze user behavior across different page categories, engagement time, and other features. Gain insights into factors influencing purchase decisions and identify areas for optimization. Formulate some hypotheses on the dataset and check if they are correct.

```
[6]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Downloading and reading the shopping csv file

```
[7]: df = pd.read_csv("shopping.csv")
df
```

```
[7]:
```

	Administrative	Administrative_Duration	Informational	\
0	0	0.0	0	
1	0	0.0	0	
2	0	0.0	0	
3	0	0.0	0	
4	0	0.0	0	
...	
12325	3	145.0	0	
12326	0	0.0	0	
12327	0	0.0	0	
12328	4	75.0	0	
12329	0	0.0	0	

	Informational_Duration	ProductRelated	ProductRelated_Duration	\
0	0.0	1	0.000000	
1	0.0	2	64.000000	
2	0.0	1	0.000000	

3	0.0	2	2.666667
4	0.0	10	627.500000
...
12325	0.0	53	1783.791667
12326	0.0	5	465.750000
12327	0.0	6	184.250000
12328	0.0	15	346.000000
12329	0.0	3	21.250000

	BounceRates	ExitRates	PageValues	SpecialDay	Month	OperatingSystems	\
0	0.200000	0.200000	0.000000	0.0	Feb		1
1	0.000000	0.100000	0.000000	0.0	Feb		2
2	0.200000	0.200000	0.000000	0.0	Feb		4
3	0.050000	0.140000	0.000000	0.0	Feb		3
4	0.020000	0.050000	0.000000	0.0	Feb		3
...
12325	0.007143	0.029031	12.241717	0.0	Dec		4
12326	0.000000	0.021333	0.000000	0.0	Nov		3
12327	0.083333	0.086667	0.000000	0.0	Nov		3
12328	0.000000	0.021053	0.000000	0.0	Nov		2
12329	0.000000	0.066667	0.000000	0.0	Nov		3

	Browser	Region	TrafficType	VisitorType	Weekend	Revenue
0	1	1	1	Returning_Visitor	False	False
1	2	1	2	Returning_Visitor	False	False
2	1	9	3	Returning_Visitor	False	False
3	2	2	4	Returning_Visitor	False	False
4	3	1	4	Returning_Visitor	True	False
...
12325	6	1	1	Returning_Visitor	True	False
12326	2	1	8	Returning_Visitor	True	False
12327	2	1	13	Returning_Visitor	True	False
12328	2	3	11	Returning_Visitor	False	False
12329	2	1	2	New_Visitor	True	False

[12330 rows x 18 columns]

```
[8]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12330 entries, 0 to 12329
Data columns (total 18 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Administrative                        12330 non-null  int64
1   Administrative_Duration              12330 non-null  float64
2   Informational                        12330 non-null  int64
```

```

3   Informational_Duration    12330 non-null float64
4   ProductRelated           12330 non-null int64
5   ProductRelated_Duration  12330 non-null float64
6   BounceRates              12330 non-null float64
7   ExitRates                12330 non-null float64
8   PageValues               12330 non-null float64
9   SpecialDay               12330 non-null float64
10  Month                    12330 non-null object
11  OperatingSystems         12330 non-null int64
12  Browser                  12330 non-null int64
13  Region                   12330 non-null int64
14  TrafficType              12330 non-null int64
15  VisitorType              12330 non-null object
16  Weekend                  12330 non-null bool
17  Revenue                  12330 non-null bool
dtypes: bool(2), float64(7), int64(7), object(2)
memory usage: 1.5+ MB

```

Unique number of values for specific categorical columns

```

[9]: columns_list = df[['SpecialDay', 'Month', 'OperatingSystems', 'Browser',
    ↪ 'Region', 'TrafficType', 'VisitorType',
    ↪ 'Weekend', 'Revenue']]

for columns in columns_list.columns:
    unique_count = columns_list[columns].nunique()
    print(columns, "-", unique_count)

```

```

SpecialDay - 6
Month - 10
OperatingSystems - 8
Browser - 13
Region - 9
TrafficType - 20
VisitorType - 3
Weekend - 2
Revenue - 2

```

Checking for the presence of null values in dataset.

```

[10]: df.isna().isna().sum()

```

```

[10]: Administrative           0
    Administrative_Duration    0
    Informational              0
    Informational_Duration     0
    ProductRelated             0
    ProductRelated_Duration    0
    BounceRates                0

```

```

ExitRates          0
PageValues         0
SpecialDay         0
Month              0
OperatingSystems   0
Browser            0
Region             0
TrafficType        0
VisitorType        0
Weekend            0
Revenue            0
dtype: int64

```

shape of the dataset

```
[11]: df.shape
```

```
[11]: (12330, 18)
```

summary statistics of the dataset

```
[12]: summary_df = df.describe()
summary_df
```

```
[12]:
```

	Administrative	Administrative_Duration	Informational	\
count	12330.000000	12330.000000	12330.000000	
mean	2.315166	80.818611	0.503569	
std	3.321784	176.779107	1.270156	
min	0.000000	0.000000	0.000000	
25%	0.000000	0.000000	0.000000	
50%	1.000000	7.500000	0.000000	
75%	4.000000	93.256250	0.000000	
max	27.000000	3398.750000	24.000000	

	Informational_Duration	ProductRelated	ProductRelated_Duration	\
count	12330.000000	12330.000000	12330.000000	
mean	34.472398	31.731468	1194.746220	
std	140.749294	44.475503	1913.669288	
min	0.000000	0.000000	0.000000	
25%	0.000000	7.000000	184.137500	
50%	0.000000	18.000000	598.936905	
75%	0.000000	38.000000	1464.157214	
max	2549.375000	705.000000	63973.522230	

	BounceRates	ExitRates	PageValues	SpecialDay	\
count	12330.000000	12330.000000	12330.000000	12330.000000	
mean	0.022191	0.043073	5.889258	0.061427	
std	0.048488	0.048597	18.568437	0.198917	

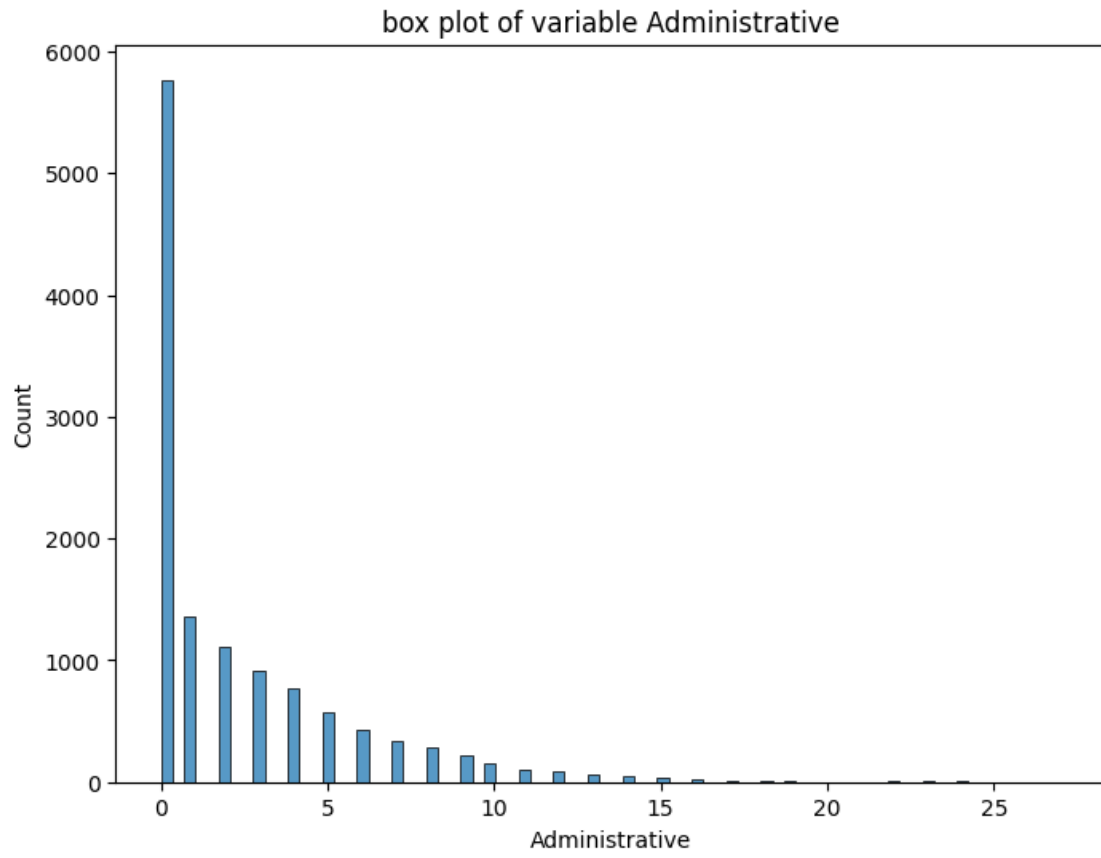
min	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.014286	0.000000	0.000000
50%	0.003112	0.025156	0.000000	0.000000
75%	0.016813	0.050000	0.000000	0.000000
max	0.200000	0.200000	361.763742	1.000000

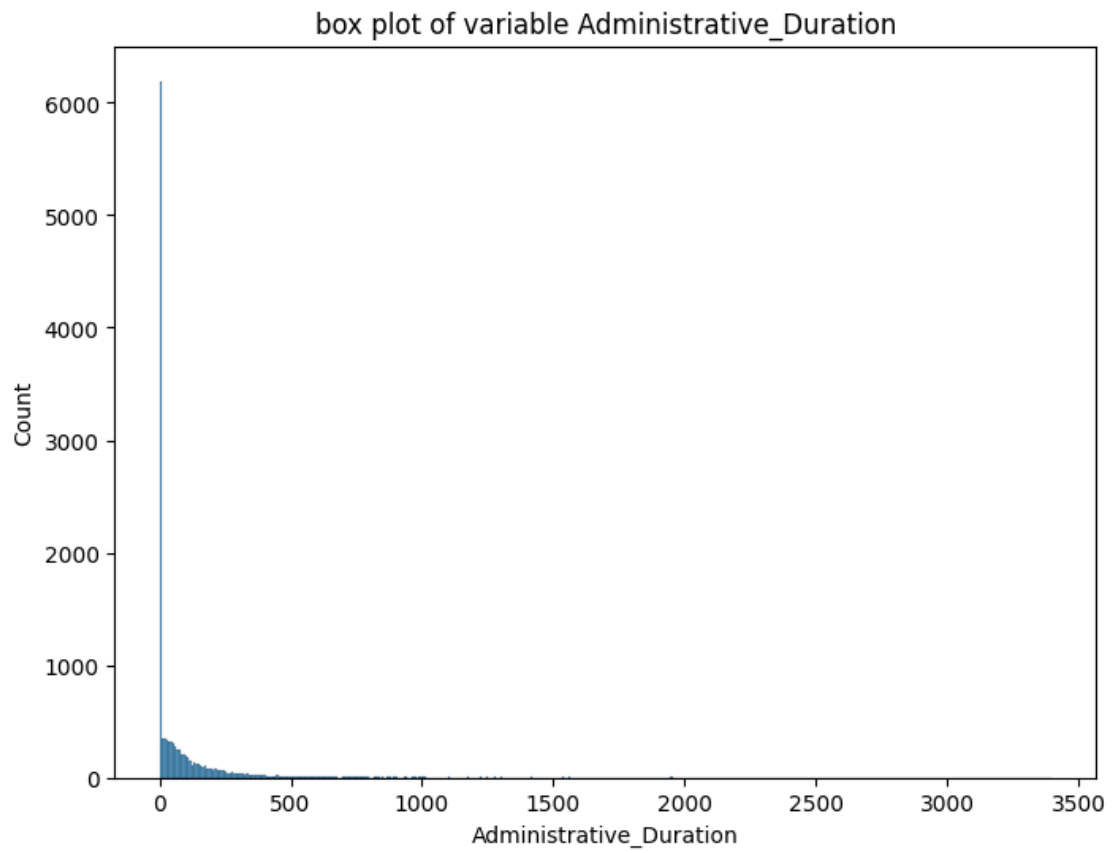
	OperatingSystems	Browser	Region	TrafficType
count	12330.000000	12330.000000	12330.000000	12330.000000
mean	2.124006	2.357097	3.147364	4.069586
std	0.911325	1.717277	2.401591	4.025169
min	1.000000	1.000000	1.000000	1.000000
25%	2.000000	2.000000	1.000000	2.000000
50%	2.000000	2.000000	3.000000	2.000000
75%	3.000000	2.000000	4.000000	4.000000
max	8.000000	13.000000	9.000000	20.000000

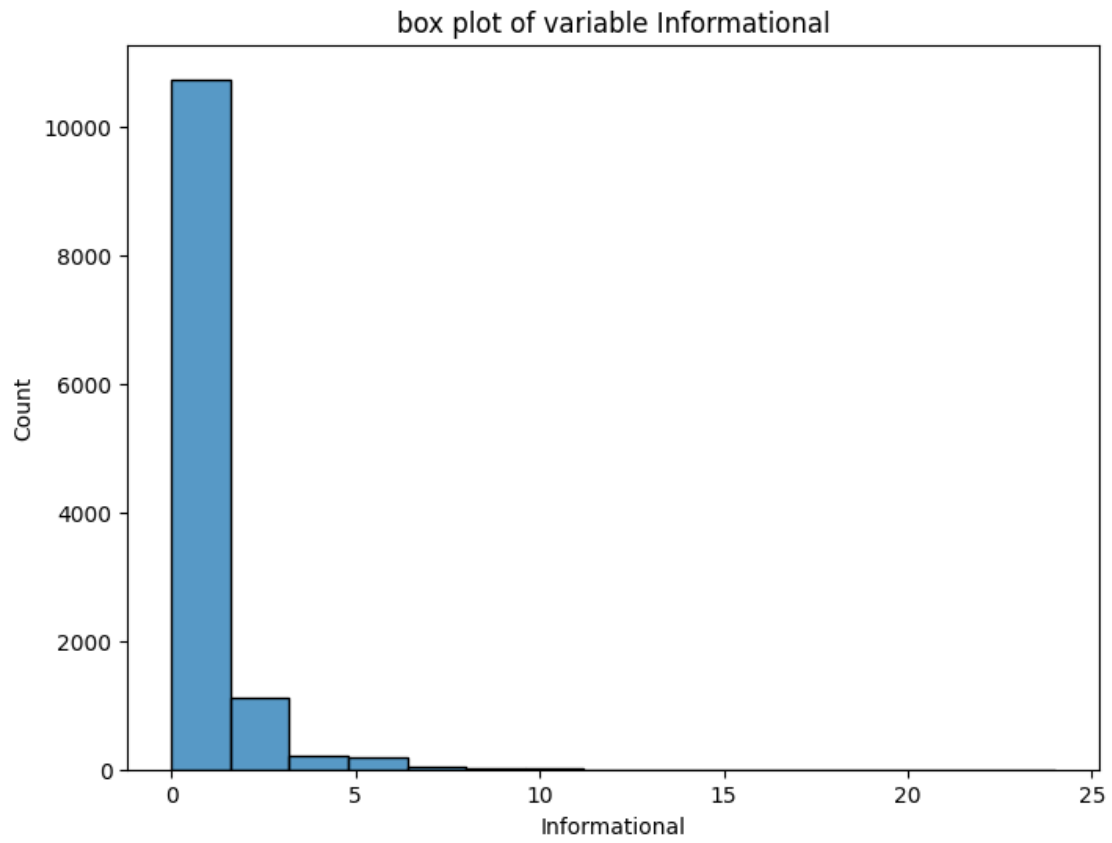
Distribution of the numerical features in the dataset

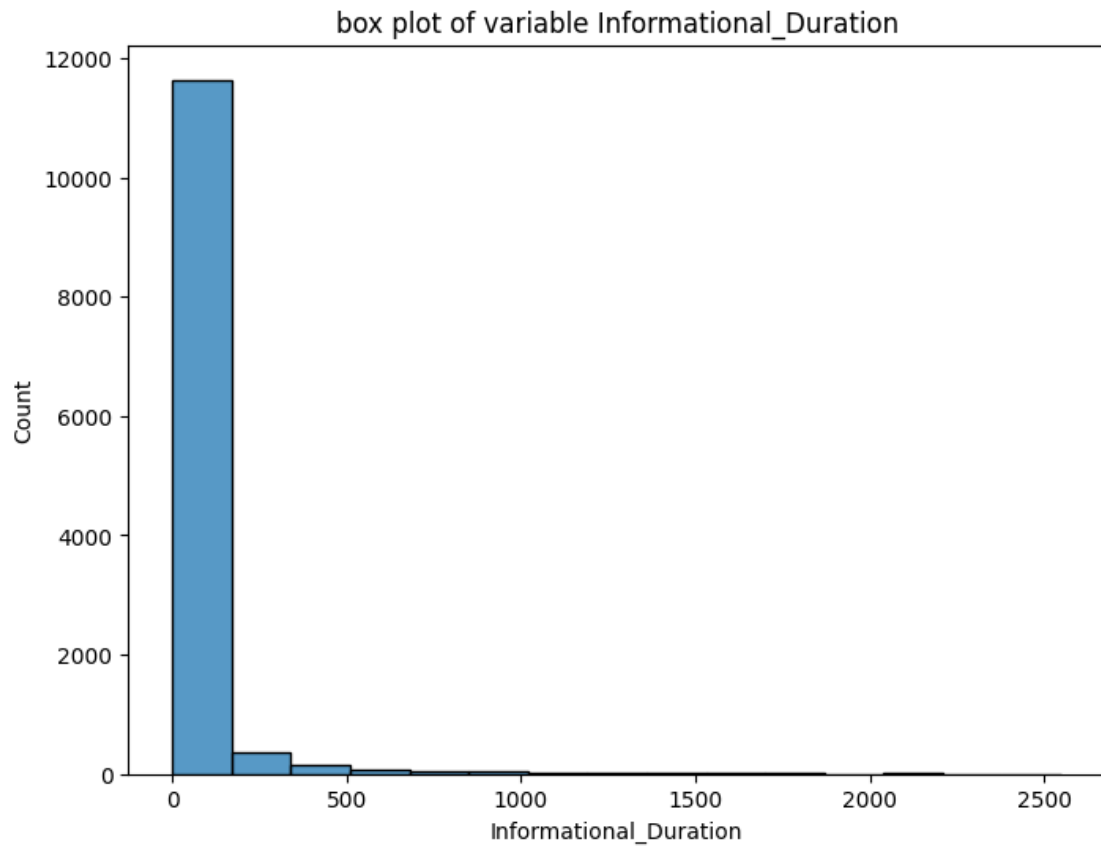
```
[13]: numeric_variables = df.select_dtypes(include = np.number).columns
      numeric_df = df[numeric_variables]

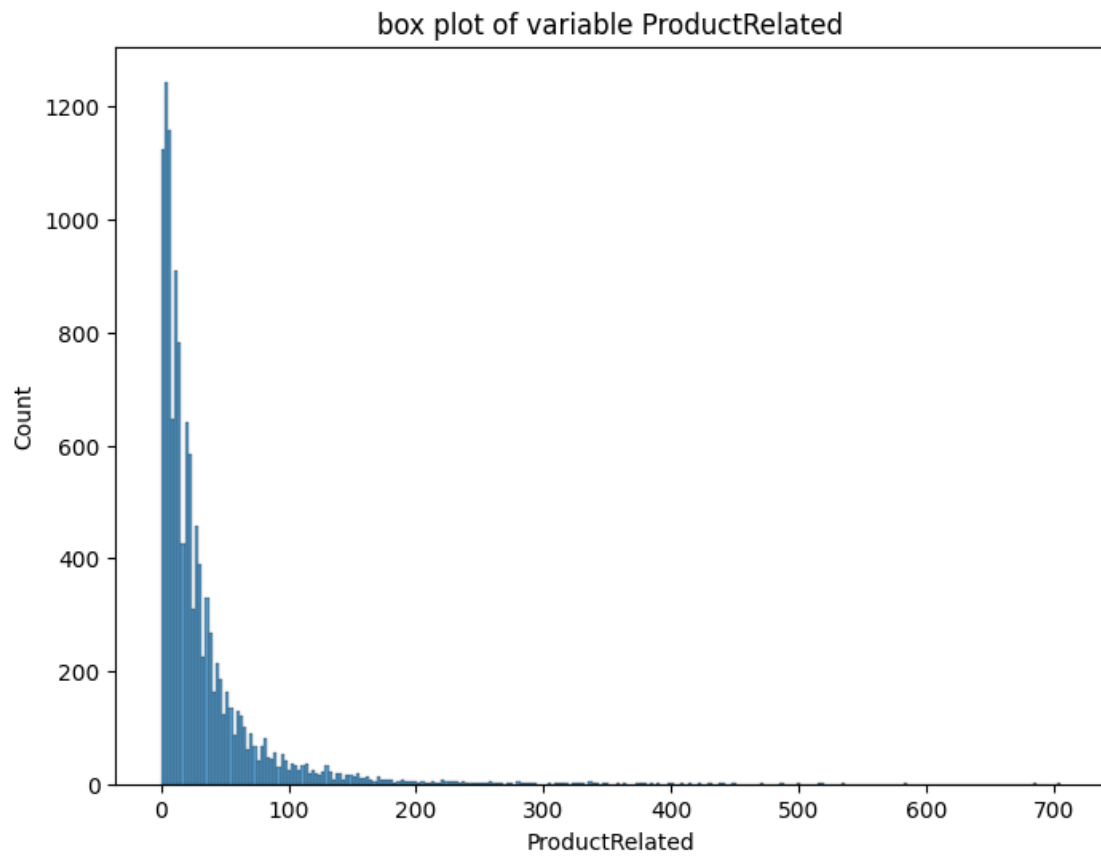
      for variable in numeric_df:
          plt.figure(figsize = (8, 6))
          sns.histplot(data = numeric_df[variable])
          plt.title(f"box plot of variable {variable}")
```

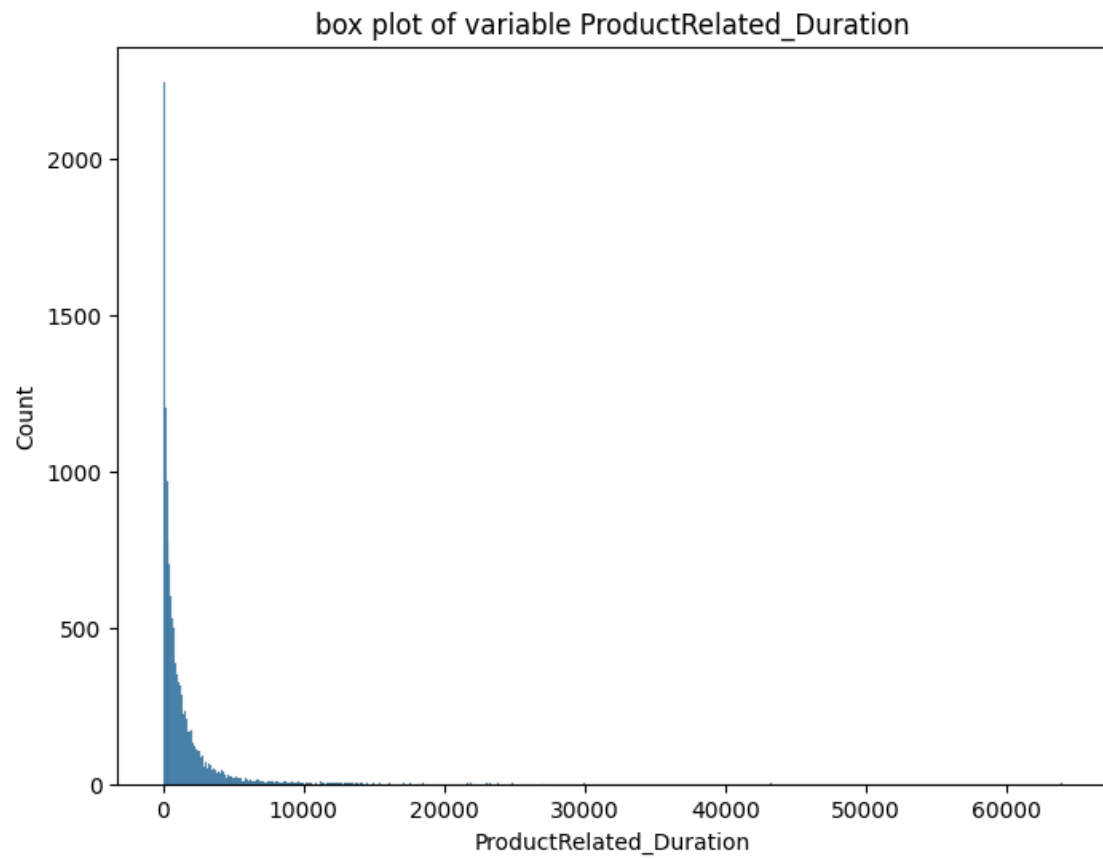





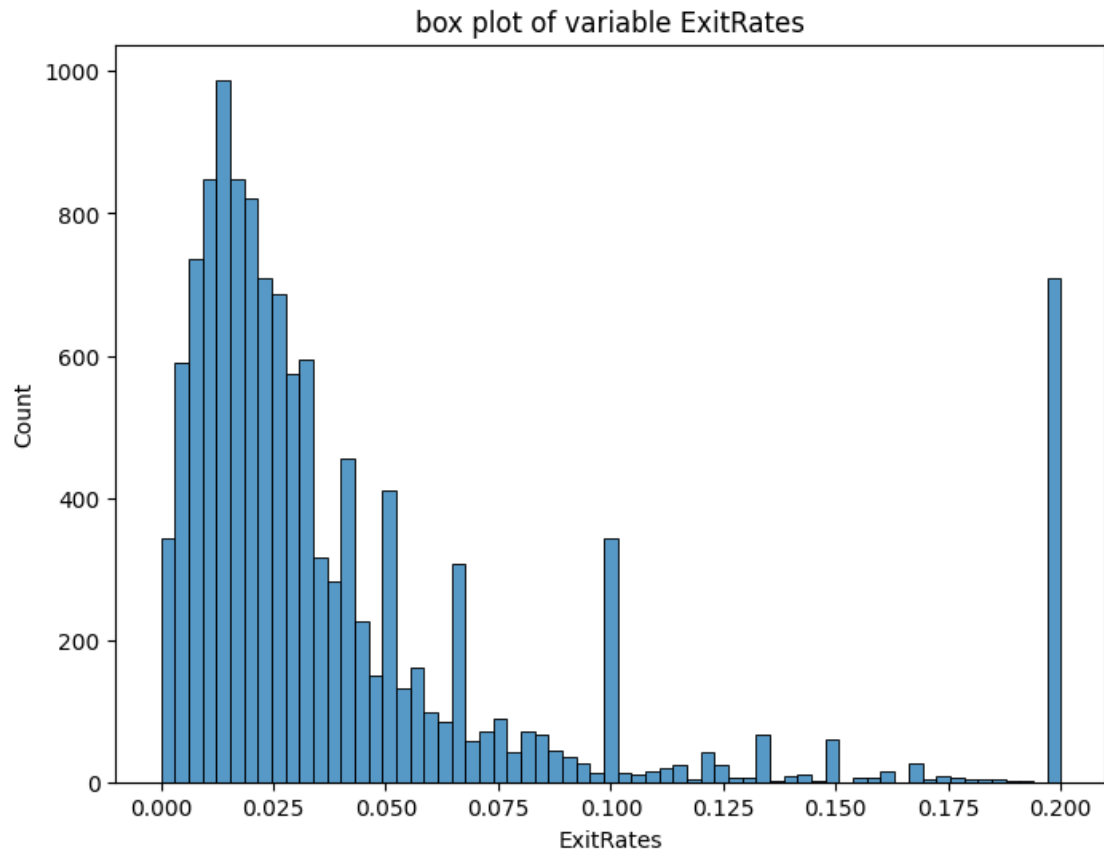


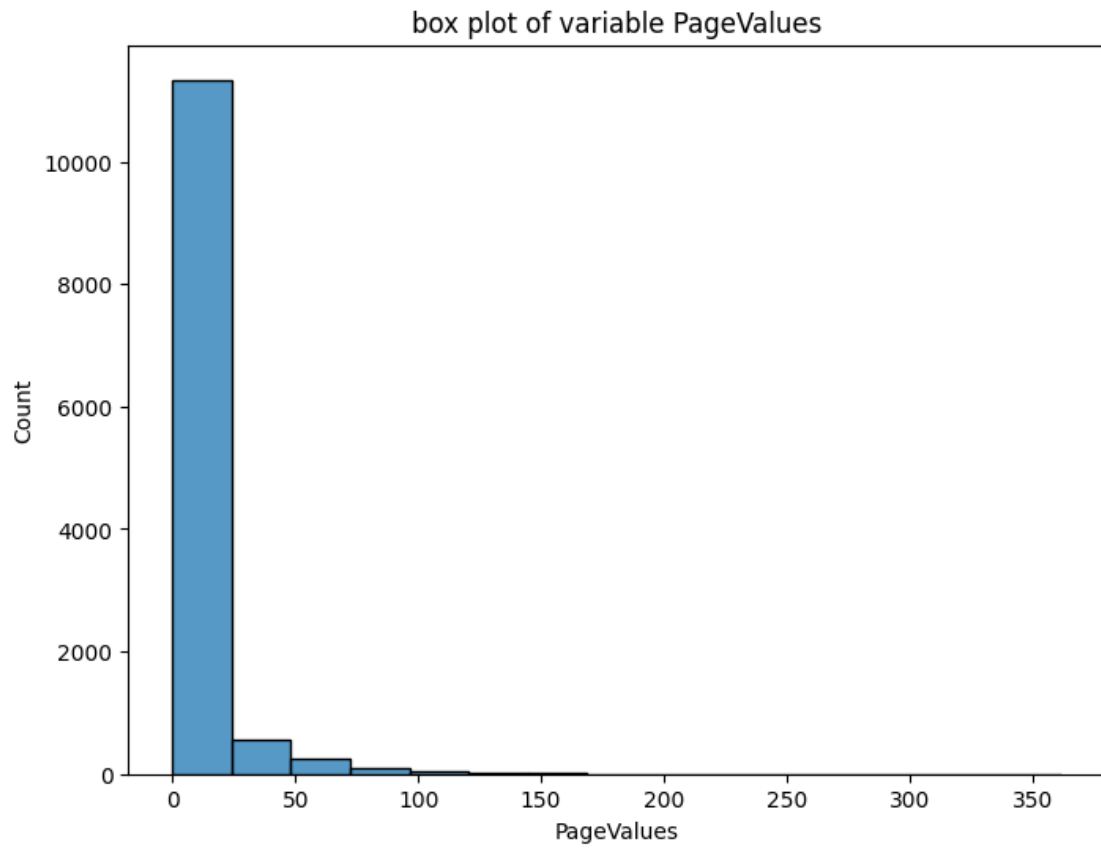


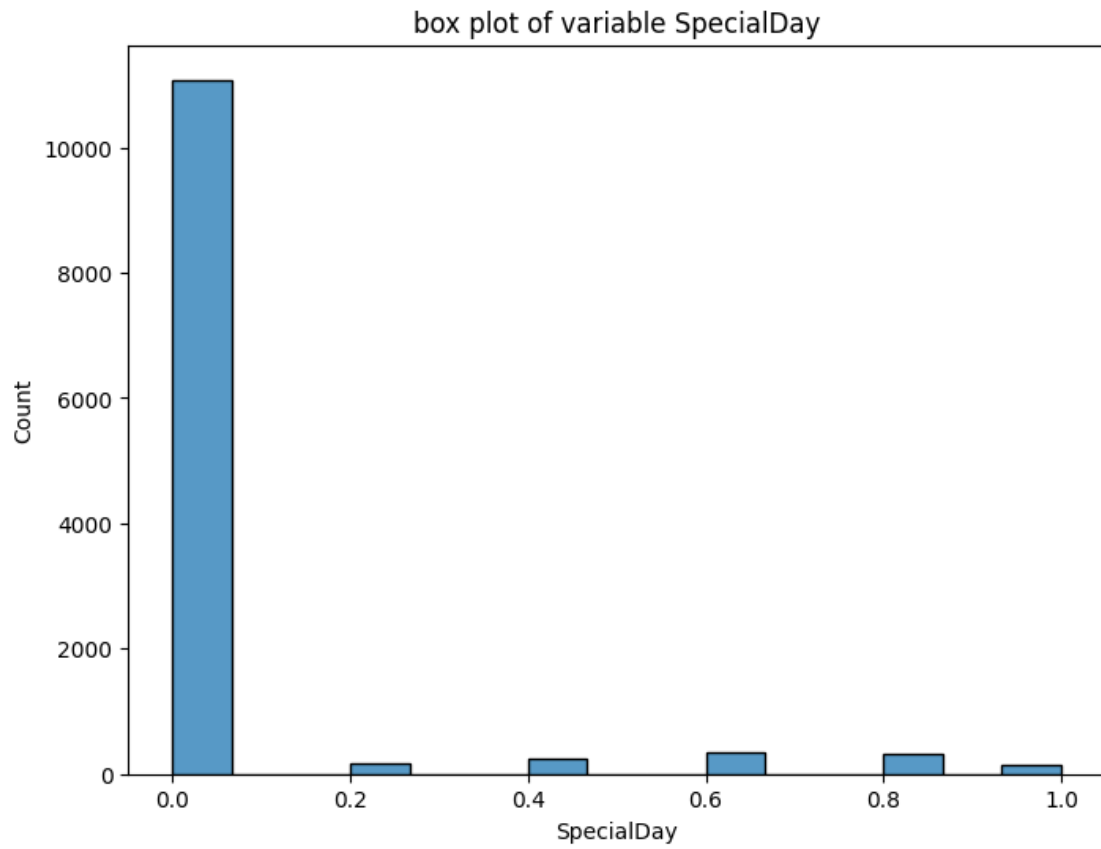


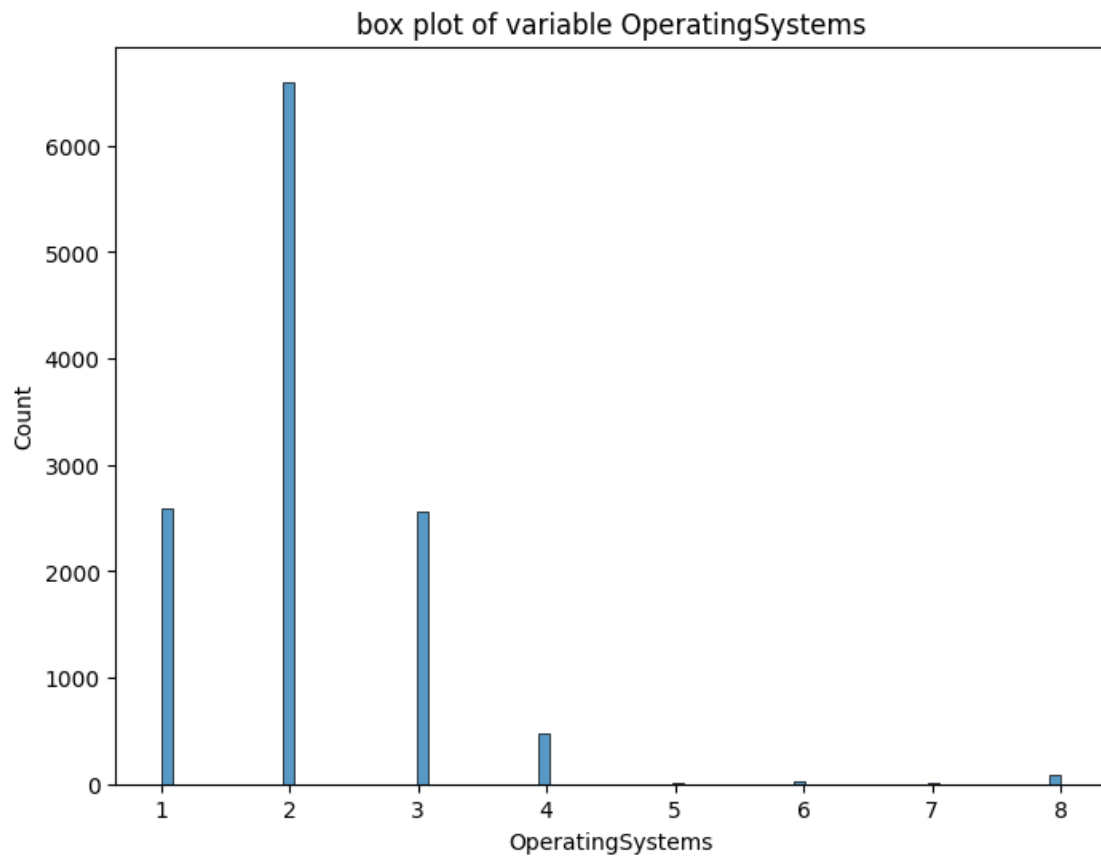


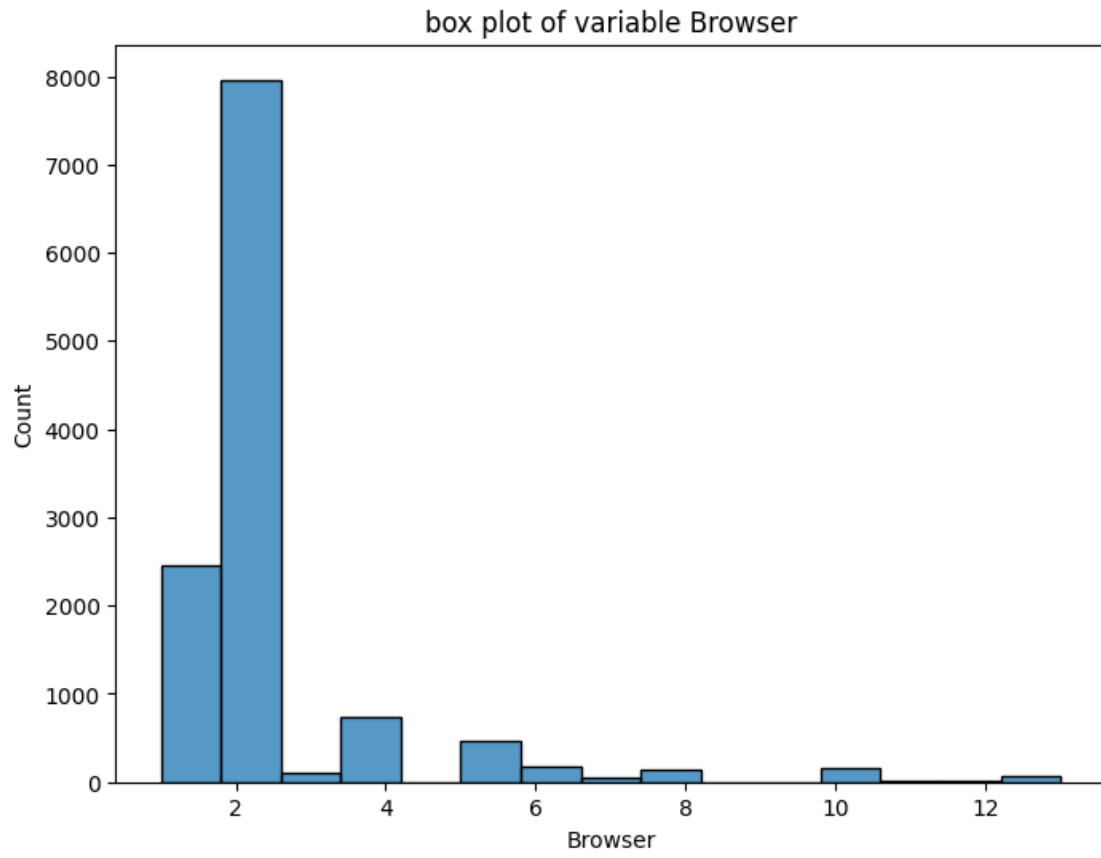


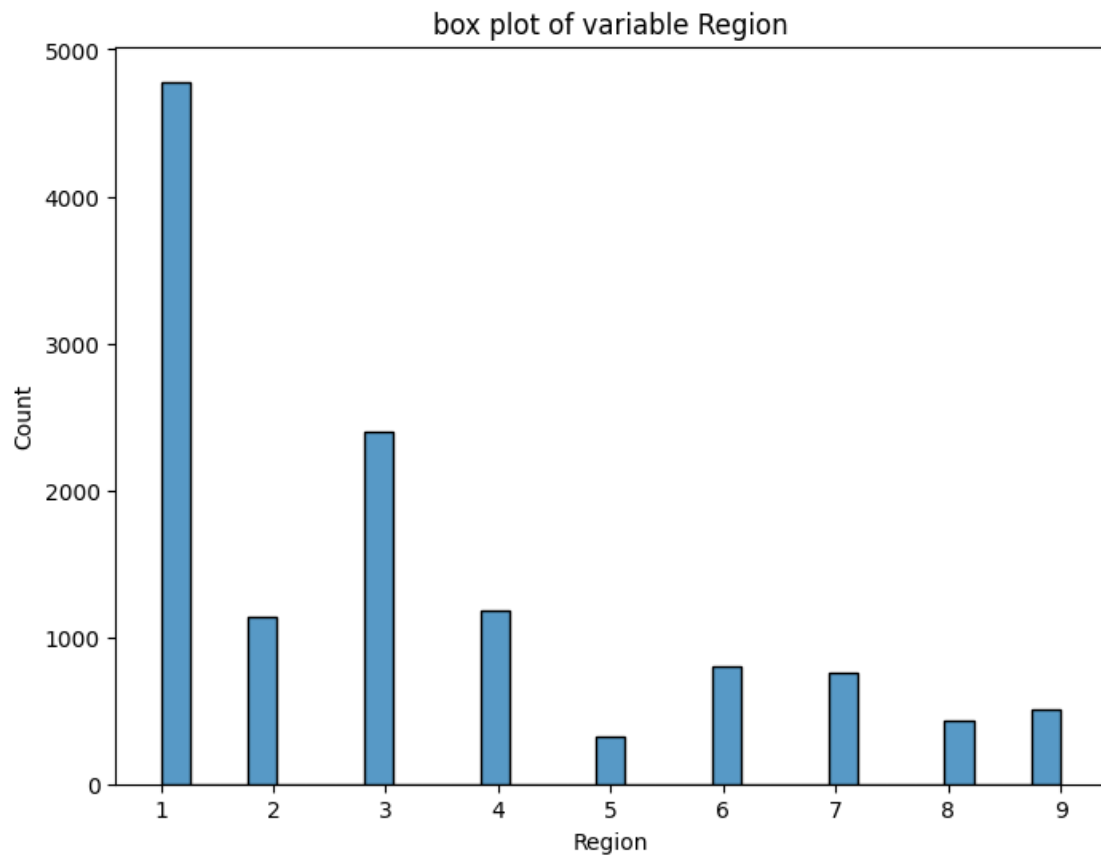


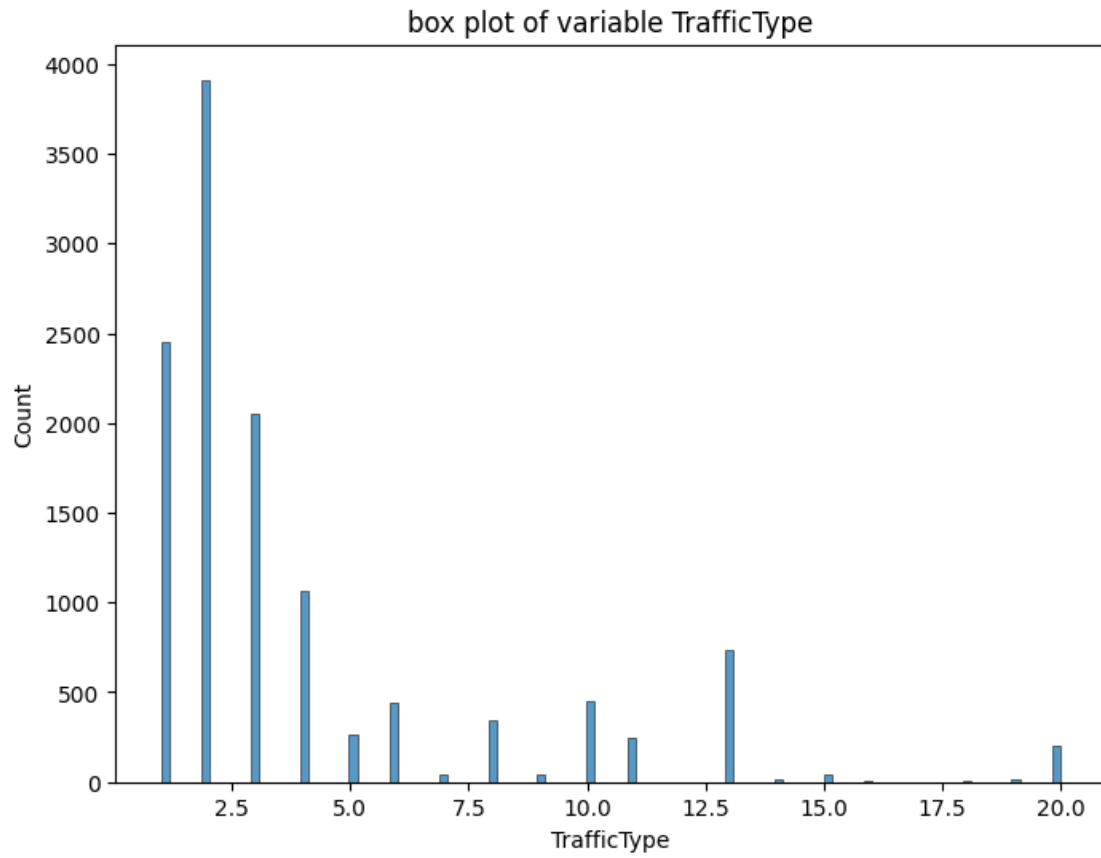











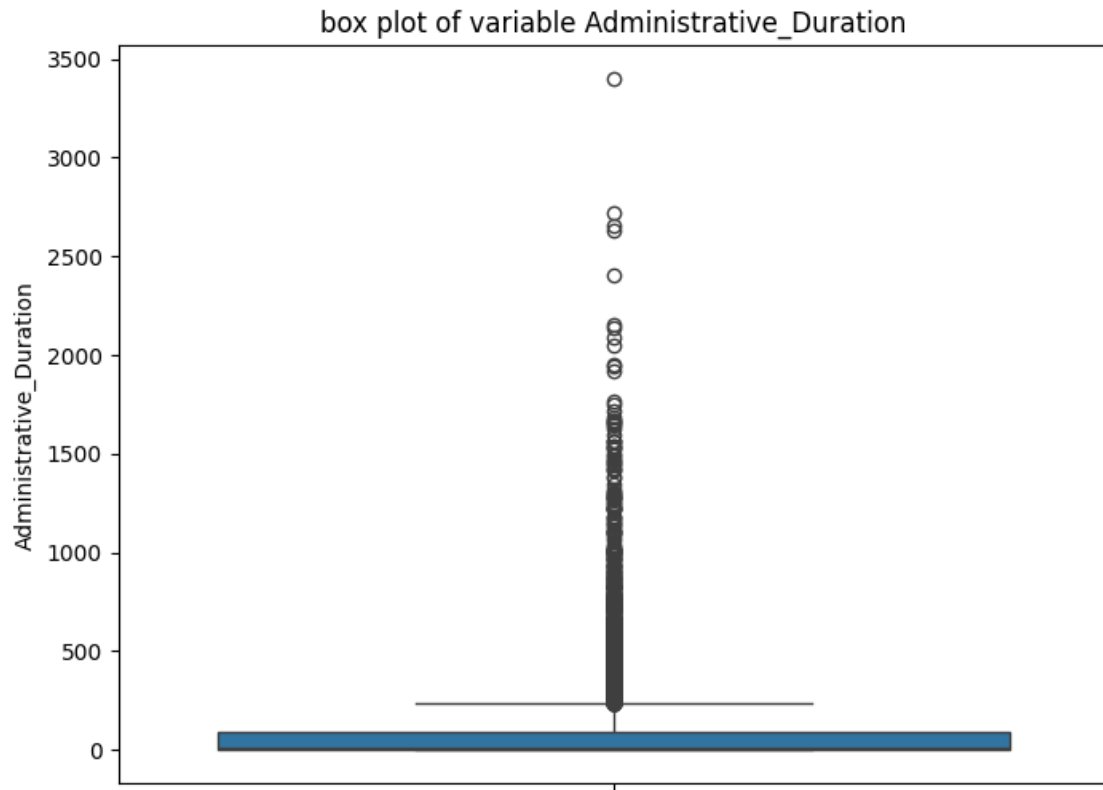


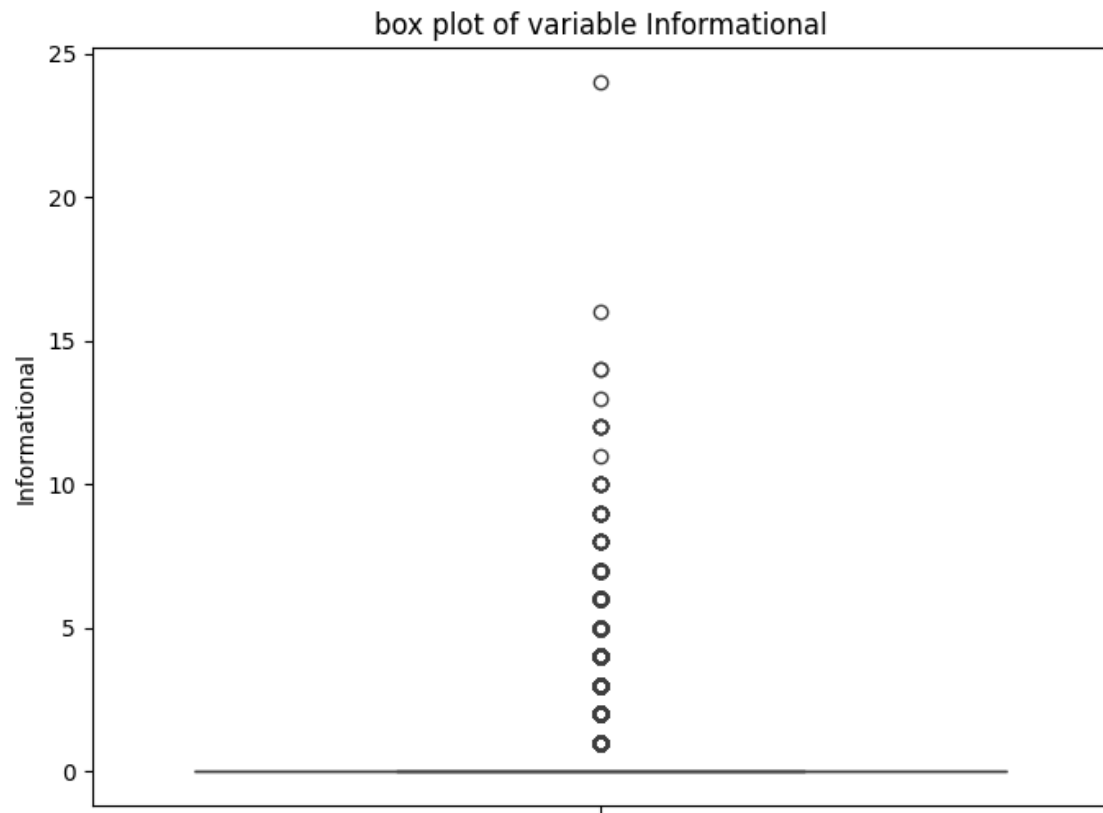
Checking for the presence of outliers in the dataset

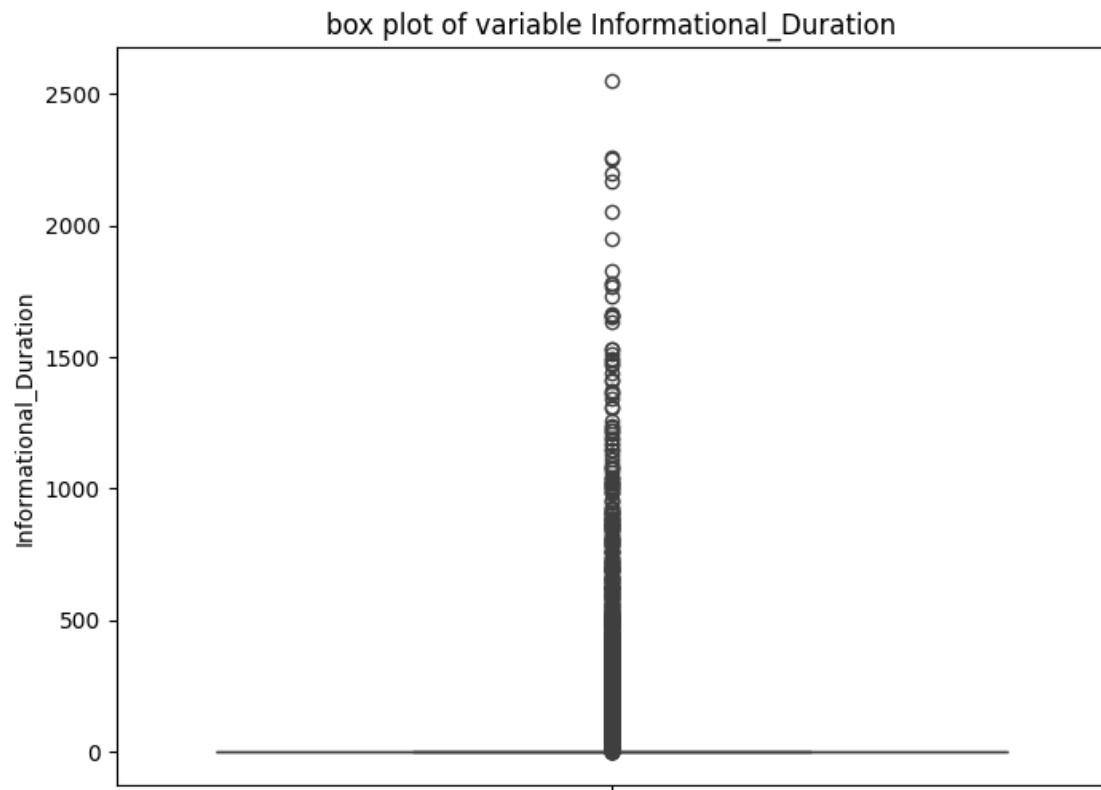
```
[42]: numeric_variables = df.select_dtypes(include = np.number).columns
      numeric_df = df[numeric_variables]

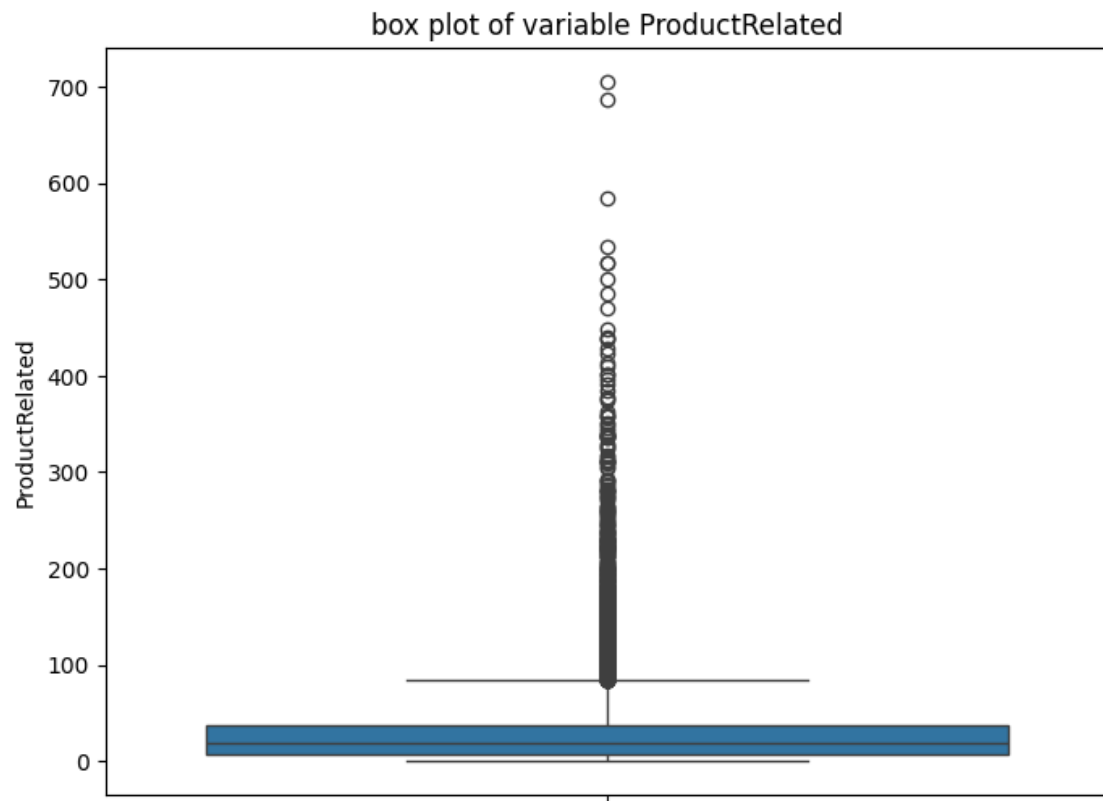
      for variable in numeric_df:
          plt.figure(figsize = (8, 6))
          sns.boxplot(data = numeric_df[variable])
          plt.title(f"box plot of variable {variable}")
```

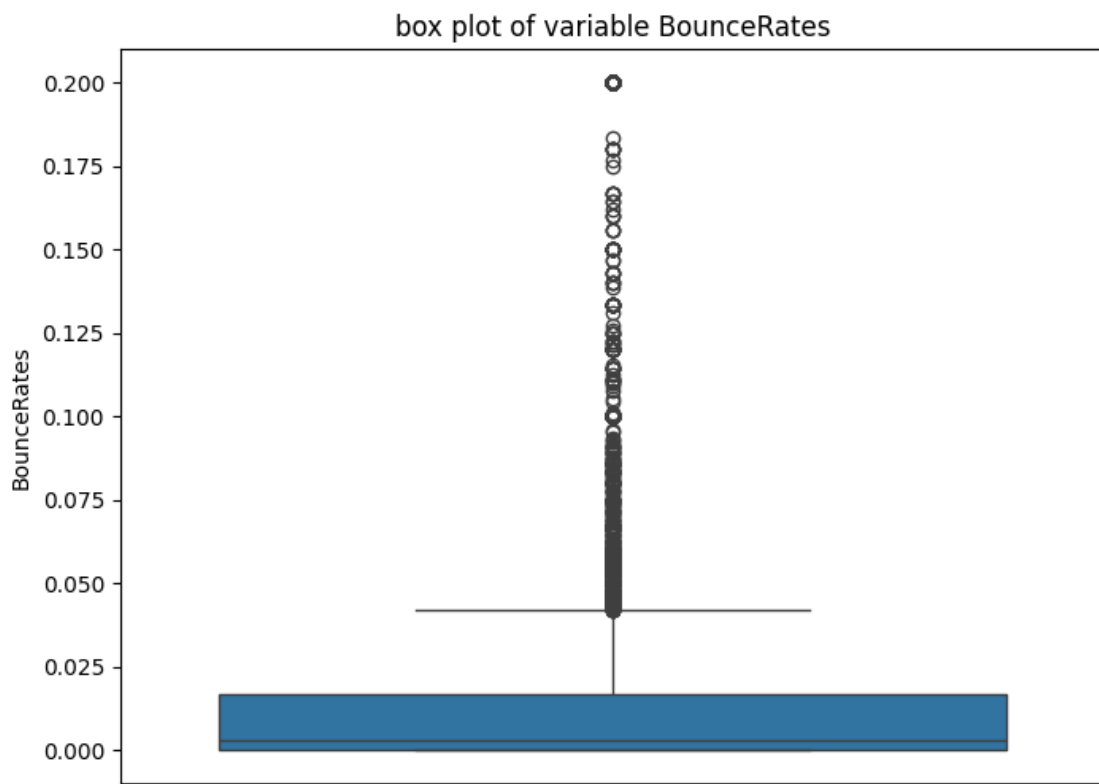



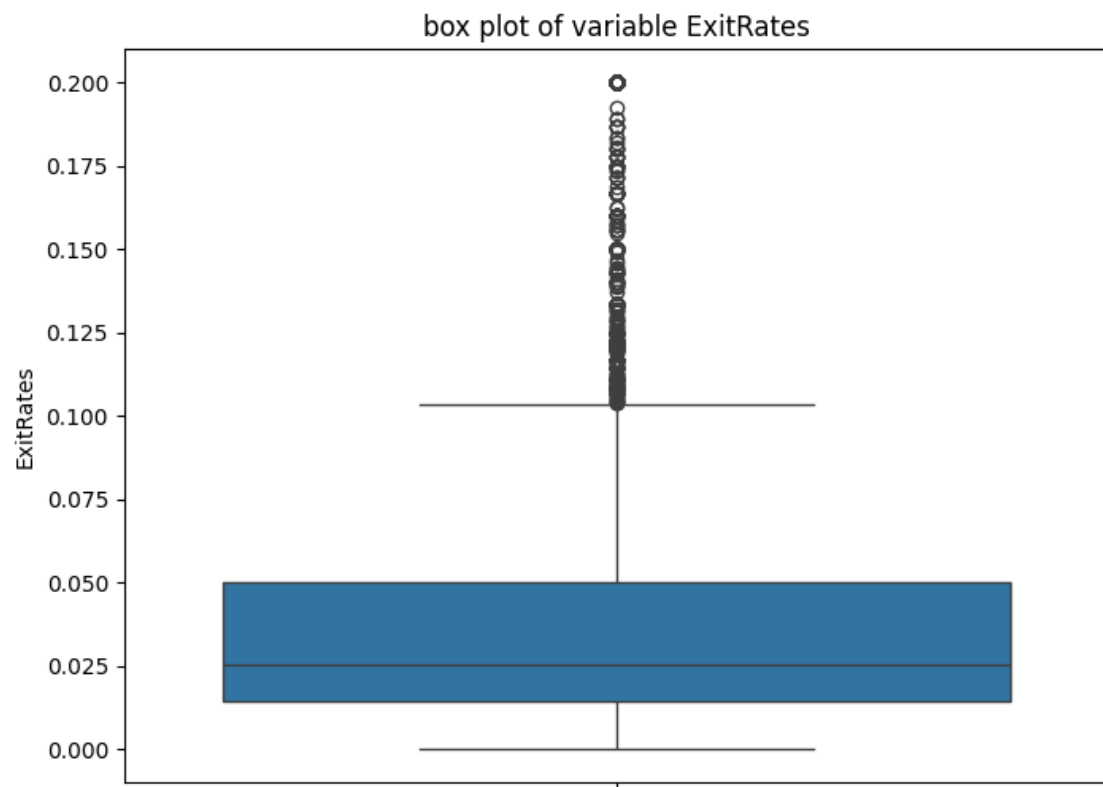














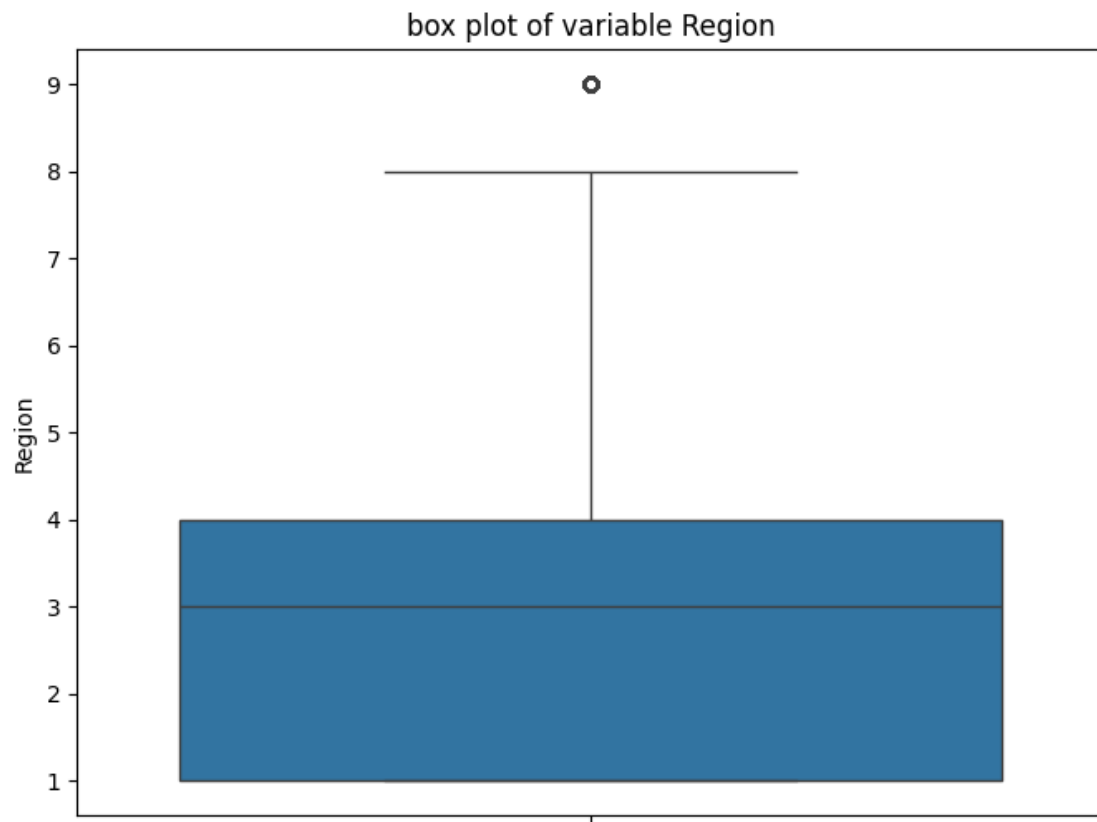















Calculating Total number of outliers

```
[9]: Q1 = summary_df[numeric_variables].loc["25%"]
      Q3 = summary_df[numeric_variables].loc["75%"]

      IQR = Q3- Q1
      print(IQR)
```

Administrative	4.000000
Administrative_Duration	93.256250
Informational	0.000000
Informational_Duration	0.000000
ProductRelated	31.000000
ProductRelated_Duration	1280.019714
BounceRates	0.016813
ExitRates	0.035714
PageValues	0.000000
SpecialDay	0.000000
OperatingSystems	1.000000
Browser	0.000000
Region	3.000000
TrafficType	2.000000

dtype: float64

```
[10]: lower_bound = Q1- 1.5*IQR
      upper_bound = Q3 + 1.5*IQR

      bounds_df = pd.DataFrame({"LowerBound" : lower_bound, "UpperBound":
      ↪upper_bound})
      print(bounds_df)
```

	LowerBound	UpperBound
Administrative	-6.000000	10.000000
Administrative_Duration	-139.884375	233.140625
Informational	0.000000	0.000000
Informational_Duration	0.000000	0.000000
ProductRelated	-39.500000	84.500000
ProductRelated_Duration	-1735.892070	3384.186784
BounceRates	-0.025219	0.042031
ExitRates	-0.039286	0.103571
PageValues	0.000000	0.000000
SpecialDay	0.000000	0.000000
OperatingSystems	0.500000	4.500000
Browser	2.000000	2.000000
Region	-3.500000	8.500000
TrafficType	-1.000000	7.000000

```
[11]: outliers_lower = (summary_df[numeric_variables] < lower_bound).sum()
      outliers_upper = (summary_df[numeric_variables] > upper_bound).sum()
      total_outliers = outliers_lower + outliers_upper

      ouliers_count_df = pd.DataFrame({"LowerBound_outliers" :outliers_lower,
      ↪"UpperBound_outliers" :outliers_upper, "Total" : total_outliers})
      print(ouliers_count_df)
```

	LowerBound_outliers	UpperBound_outliers	Total
Administrative	0	2	2
Administrative_Duration	0	2	2
Informational	0	4	4
Informational_Duration	0	4	4
ProductRelated	0	2	2
ProductRelated_Duration	0	2	2
BounceRates	0	3	3
ExitRates	0	2	2
PageValues	0	4	4
SpecialDay	0	4	4
OperatingSystems	0	2	2
Browser	2	3	5
Region	0	2	2
TrafficType	0	2	2

Understanding correlation between multiple numerical features

```
[14]: variables_1 = numeric_df[['Administrative', 'Administrative_Duration',
    ↳ 'Informational', 'Informational_Duration', 'ProductRelated',
    ↳ 'ProductRelated_Duration', 'BounceRates', 'ExitRates']]
spearman_corr = variables_1.corr(method = 'spearman')
spearman_corr
```

```
[14]:
```

	Administrative	Administrative_Duration \
Administrative	1.000000	0.940725
Administrative_Duration	0.940725	1.000000
Informational	0.369194	0.357150
Informational_Duration	0.362861	0.352060
ProductRelated	0.460204	0.430072
ProductRelated_Duration	0.421613	0.413765
BounceRates	-0.155219	-0.163609
ExitRates	-0.434389	-0.437912

	Informational	Informational_Duration \
Administrative	0.369194	0.362861
Administrative_Duration	0.357150	0.352060
Informational	1.000000	0.950958
Informational_Duration	0.950958	1.000000
ProductRelated	0.368673	0.361032
ProductRelated_Duration	0.367522	0.362720
BounceRates	0.005753	-0.002474
ExitRates	-0.185691	-0.200056

	ProductRelated	ProductRelated_Duration	BounceRates \
Administrative	0.460204	0.421613	-0.155219
Administrative_Duration	0.430072	0.413765	-0.163609
Informational	0.368673	0.367522	0.005753
Informational_Duration	0.361032	0.362720	-0.002474
ProductRelated	1.000000	0.882672	-0.052305
ProductRelated_Duration	0.882672	1.000000	-0.079768
BounceRates	-0.052305	-0.079768	1.000000
ExitRates	-0.518920	-0.476935	0.602276

	ExitRates
Administrative	-0.434389
Administrative_Duration	-0.437912
Informational	-0.185691
Informational_Duration	-0.200056
ProductRelated	-0.518920
ProductRelated_Duration	-0.476935
BounceRates	0.602276
ExitRates	1.000000

Understanding correlation between Pages visits and Page values

```
[18]: variables_2 = numeric_df[['Administrative', 'Informational', 'ProductRelated', 'PageValues']]
      spearman_corr = variables_2.corr(method = 'spearman')
      spearman_corr
```

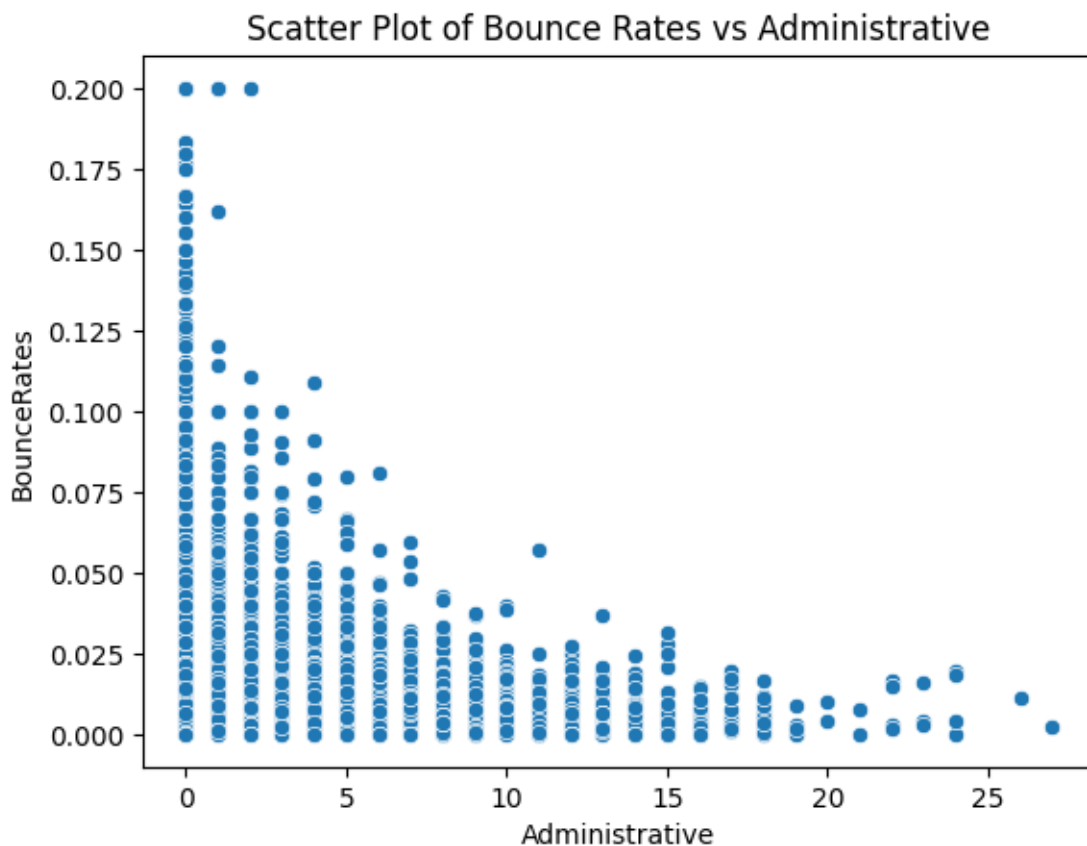
```
[18]:
```

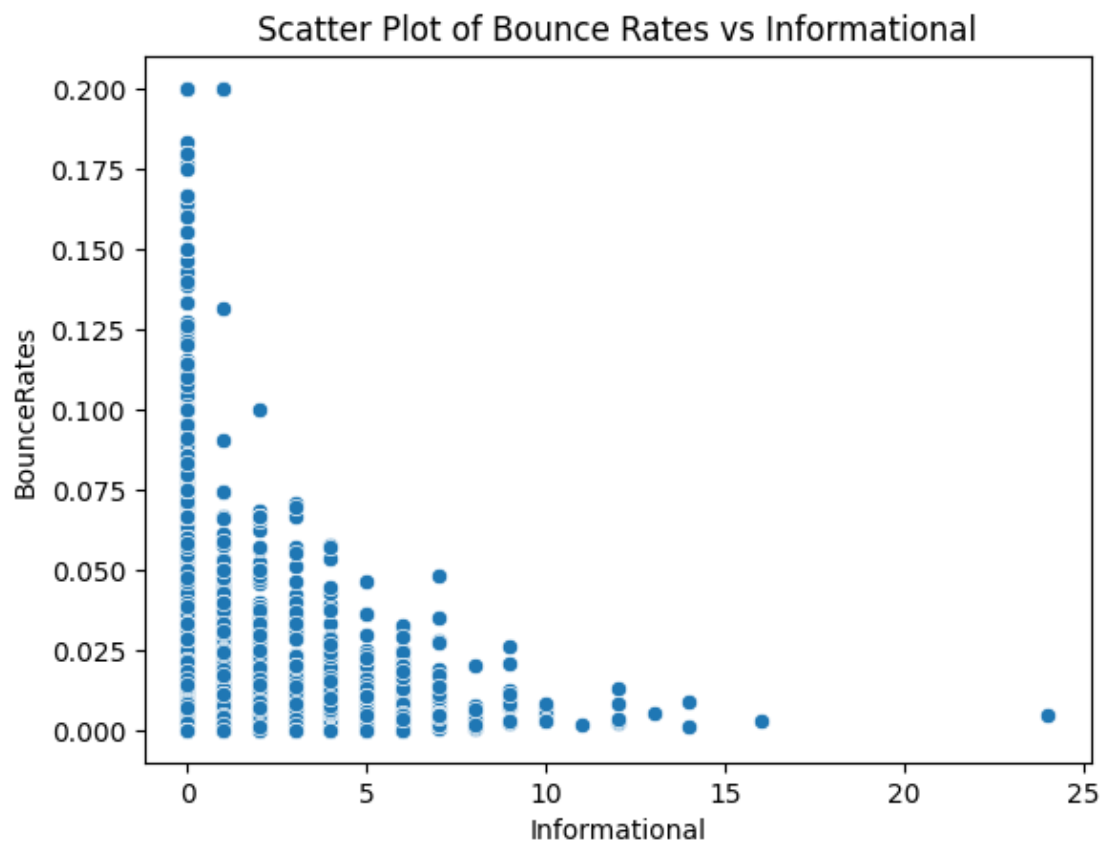
	Administrative	Informational	ProductRelated	PageValues
Administrative	1.000000	0.369194	0.460204	0.328350
Informational	0.369194	1.000000	0.368673	0.219471
ProductRelated	0.460204	0.368673	1.000000	0.341975
PageValues	0.328350	0.219471	0.341975	1.000000

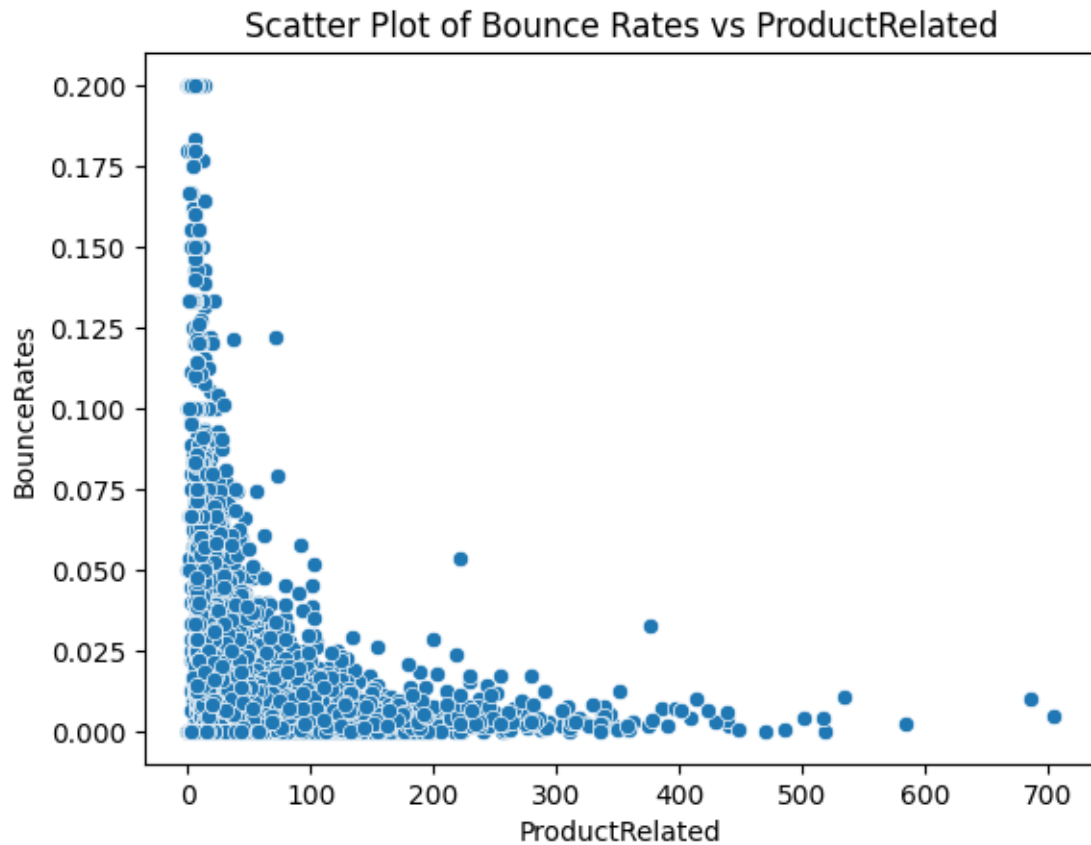
Scatter plot showing relationship between Pages visited and bounce rate

```
[23]: columns_to_plot = ['Administrative', 'Informational', 'ProductRelated']

for col in columns_to_plot:
    sns.scatterplot(x=df[col], y=df['BounceRates'])
    plt.title(f'Scatter Plot of Bounce Rates vs {col}')
    plt.show()
```









Scatter plot showing relationship between Pages visited and Exit rate


```
[24]: columns_to_plot = ['Administrative', 'Informational', 'ProductRelated']

for col in columns_to_plot:
    sns.scatterplot(x=df[col], y=df['ExitRates'])
    plt.title(f'Scatter Plot of Exit Rates vs {col}')
    plt.show()
```

Scatter Plot of Exit Rates vs Administrative







Total pages visited & Individual Revenue contribution for each type of Operating system

```
[24]: pages_under_OS = df.groupby('OperatingSystems')[['Administrative',
↳ 'Informational', 'ProductRelated', 'Revenue']].sum()
pages_under_OS = pages_under_OS.sort_values(by = ['Administrative',
↳ 'Informational', 'ProductRelated'], ascending = False).reset_index()
pages_under_OS
```

```
[24]:
```

	OperatingSystems	Administrative	Informational	ProductRelated	Revenue
0	2	15620	3329	243291	1155
1	3	5930	1402	67616	268
2	1	5756	1275	61224	379
3	4	1049	171	17317	85
4	8	123	17	1182	17
5	6	39	10	406	2
6	7	25	4	140	1
7	5	4	1	73	1

Total pages visited & Individual Revenue contribution for each type of Web browser


```
[29]: pages_under_browser = df.groupby('Browser')[['Administrative', 'Informational',
↪ 'ProductRelated', 'Revenue']].sum()
pages_under_browser = pages_under_browser.sort_values(by = ['Administrative',
↪ 'Informational', 'ProductRelated'], ascending = False).reset_index()
pages_under_browser
```

```
[29]:
```

	Browser	Administrative	Informational	ProductRelated	Revenue
0	2	19345	4327	276985	1223
1	1	5376	1200	60086	365
2	4	1460	268	22411	130
3	5	1056	183	14633	86
4	10	376	72	5357	32
5	6	334	72	5198	20
6	8	189	19	2433	21
7	3	141	28	1724	5
8	7	137	23	1259	6
9	13	107	14	908	16
10	12	19	2	172	3
11	11	4	1	73	1
12	9	2	0	10	0

Total pages visited & Individual Revenue contribution for each type of Region

```
[30]: pages_under_Region = df.groupby('Region')[['Administrative', 'Informational',
↪ 'ProductRelated', 'Revenue']].sum()
pages_under_Region = pages_under_Region.sort_values(by = ['Administrative',
↪ 'Informational', 'ProductRelated'], ascending = False).reset_index()
print(pages_under_Region)

sns.barplot(data = pages_under_Region.melt(id_vars = 'Region'), x = 'Region', y_
↪ = 'value', hue = 'variable', color = 'yellow')
```


	Region	Administrative	Informational	ProductRelated	Revenue
0	1	10857	2607	161567	771
1	3	5878	1228	77102	349
2	2	2731	506	34566	188
3	4	2722	604	35706	175
4	7	1792	339	23941	119
5	6	1655	408	22410	112
6	8	1077	196	11600	56
7	9	1051	183	14935	86
8	5	783	138	9422	52

<ipython-input-30-49afc0e1f0f6>:5: FutureWarning:

Setting a gradient palette using color= is deprecated and will be removed in v0.14.0. Set `palette='dark:yellow'` for the same effect.

```
sns.barplot(data = pages_under_Region.melt(id_vars = 'Region'), x = 'Region',
y = 'value', hue = 'variable', color = 'yellow')
```

```
[30]: <Axes: xlabel='Region', ylabel='value'>
```



Calculation of weekend and weekdays proportions

```
[25]: df['Weekend'].value_counts(normalize = True).reset_index()
```

```
[25]:   Weekend  proportion
0    False    0.767397
1     True    0.232603
```

Checking Total revenue class balance

```
[27]: df['Revenue'].value_counts(normalize = True).reset_index()
```

```
[27]:   Revenue  proportion
0    False    0.845255
1     True    0.154745
```

Checking for which type of visitors actually contributed more to revenue

```
[37]: counts = df.groupby('VisitorType')['Revenue'].value_counts().reset_index()
counts
```

```
[37]:
```

	VisitorType	Revenue	count
0	New_Visitor	False	1272
1	New_Visitor	True	422
2	Other	False	69
3	Other	True	16
4	Returning_Visitor	False	9081
5	Returning_Visitor	True	1470

Total pages visited & Individual Revenue contribution for each type of Month

```
[36]: Monthly_page_views = df.groupby('Month')[['Administrative', 'Informational', 'ProductRelated', 'Revenue']].sum()
Monthly_page_views = Monthly_page_views.sort_values(by = ['Administrative', 'Informational', 'ProductRelated'], ascending = False).reset_index()
Monthly_page_views
```

```
[36]:
```

	Month	Administrative	Informational	ProductRelated	Revenue
0	Nov	7847	1938	138024	760
1	May	6610	1426	89105	365
2	Dec	3793	885	48347	216
3	Mar	3600	802	37775	192
4	Oct	2042	268	18428	115
5	Sep	1494	254	14831	86
6	Aug	1358	235	16566	76
7	Jul	1047	223	15728	66
8	June	655	162	10387	29
9	Feb	100	16	2058	3

Calculation of total sessions where a user visited all types of pages

```
[15]: df['visited_all'] = (df['Administrative'] > 0) & (df['Informational'] > 0) & (df['ProductRelated'] > 0)
df['visited_all'].sum()
```

```
[15]: 2167
```

Calculation of total sessions where a user visited all types of pages and also contributed to revenue

```
[16]: df['visited_all&Purchased'] = (df['Administrative'] > 0) & (df['Informational'] > 0) & (df['ProductRelated'] > 0) & (df['Revenue'] > 0)
df['visited_all&Purchased'].sum()
```

```
[16]: 526
```

Finding correlation b/w Special day and revenue

```
[17]: variables_3 = df[['SpecialDay', 'Revenue']]
spearman_corr = variables_3.corr(method = 'spearman')
spearman_corr
```

```
[17]:          SpecialDay  Revenue
SpecialDay    1.000000 -0.086878
Revenue      -0.086878  1.000000
```


Investigating user session lengths and their impact on conversion rates.

```
[18]: df['Total_time_spent'] = df['Administrative_Duration'] +
    ↪df['Informational_Duration'] + df['ProductRelated_Duration']

variables_4 = df[['Total_time_spent', 'Revenue']]
spearman_corr = variables_4.corr(method = 'spearman')
print(spearman_corr)

sns.histplot(data = df, x = 'Total_time_spent', y = 'Revenue', kde = True,
    ↪color = 'red')
plt.title('Histogram of Session Length by Purchases')
plt.show()
```

```
          Total_time_spent  Revenue
Total_time_spent    1.000000  0.220721
Revenue             0.220721  1.000000
```



Exploring PageValues distribution and its relationship with TrafficType, VisitorType, and Region.

```
[19]: variables_5 = df[['TrafficType', 'VisitorType', 'Region']]

for variable in variables_5:
    grouped_data = df.groupby(variable)['PageValues'].sum()
    print(grouped_data)
    print("<----->")


    sns.lineplot(x = grouped_data.index, y = grouped_data.values)
    plt.title(f'Total Page Values by {variable}')
    plt.xlabel(variable)
    plt.ylabel('Total page views')
    plt.show()
```

```
TrafficType
1      8468.386672
2     32494.983720
3      6722.420075
4      7529.087303
```

5	2005.247088
6	2253.852296
7	542.693810
8	3533.735395
9	160.379441
10	2793.703633
11	1251.954486
12	0.000000
13	1737.684088
14	64.169261
15	1.385792
16	0.000000
17	0.000000
18	0.000000
19	59.457846
20	2995.408541

Name: PageValues, dtype: float64


<----->



VisitorType	
New_Visitor	18248.085596

```
Other          1546.304039
Returning_Visitor  52820.159812
Name: PageValues, dtype: float64
```


<----->



```
Region
1      28212.256912
2       6709.893181
3      12842.429647
4       6860.170650
5       2941.753454
6       3866.874528
7       4763.541048
8       1847.940453
9       4569.689575
```

```
Name: PageValues, dtype: float64
```

<----->



Grouping users based on VisitorType, OperatingSystems, and Region to identify potential differences in behavior and conversion rates


```
[73]: variables_6 = df[['VisitorType', 'OperatingSystems', 'Region']]

for variable in variables_6:
    grouped_data = df.groupby(variable)['Revenue'].sum()
    print(grouped_data)
    print("<----->")

    sns.barplot(x = grouped_data.index, y = grouped_data.values, color = 'green')
    plt.title(f'Total Purchases by {variable}')
    plt.xlabel(variable)
    plt.ylabel('Purchase contribution')
    plt.show()
```

```
VisitorType
New_Visitor      422
Other             16
Returning_Visitor 1470
Name: Revenue, dtype: int64
```


<----->




OperatingSystems

1	379
2	1155
3	268
4	85
5	1
6	2
7	1
8	17

Name: Revenue, dtype: int64

<----->




Region

1	771
2	188
3	349
4	175
5	52
6	112
7	119
8	56
9	86

Name: Revenue, dtype: int64

<----->



Hypothesis testing: Performing One- way and Two way Anova to check relationship between Independent and dependent variable.

```
[20]: import statsmodels.api as sm
      from statsmodels.formula.api import ols
```

```
[21]: df['Revenue'] = df['Revenue'].replace({True : 1, False: 0})
      df['visited_all'] = df['visited_all'].replace({True : 1, False : 0})
      df['visited_all&Purchased'] = df['visited_all&Purchased'].replace({True: 1,
      ↪False : 0})
      df['Weekend'] = df['Weekend'].replace({True : 1, False : 0})
```

```
[23]: # H0: There is no relationship between dependent and independent variables.
      # H1: There is a signifiant relationship between dependent and independent
      ↪variables.

      test = ols('Revenue ~ C(OperatingSystems) * C(Browser)', data =df).fit()
      anova_table = sm.stats.anova_lm(test,typ= 1)
      print(anova_table)
```

	df	sum_sq	mean_sq	F	\
C(OperatingSystems)	7.0	9.813438	1.401920	10.801121	
C(Browser)	12.0	1.867781	0.155648	1.199197	
C(OperatingSystems):C(Browser)	84.0	8.962134	0.106692	0.822011	
Residual	12288.0	1594.907520	0.129794	NaN	

	PR(>F)
C(OperatingSystems)	1.194180e-13
C(Browser)	2.766328e-01
C(OperatingSystems):C(Browser)	8.799230e-01
Residual	NaN

```
[24]: test = ols('Revenue ~ C(visited_all)', data =df).fit()
anova_table = sm.stats.anova_lm(test,typ= 1)
print(anova_table)
```

	df	sum_sq	mean_sq	F	PR(>F)
C(visited_all)	1.0	20.353573	20.353573	157.573355	6.359801e-36
Residual	12328.0	1592.393872	0.129169	NaN	NaN

```
[25]: test = ols('Revenue ~ C(SpecialDay)', data =df).fit()
anova_table = sm.stats.anova_lm(test,typ= 1)
print(anova_table)
```

	df	sum_sq	mean_sq	F	PR(>F)
C(SpecialDay)	5.0	12.566730	2.513346	19.356862	3.003311e-19
Residual	12324.0	1600.180715	0.129843	NaN	NaN

```
[1]: from google.colab import drive
drive.mount("/content/drive")
```

Mounted at /content/drive

```
[ ]: !pip install nbconvert
!apt-get install texlive texlive-xetex texlive-latex-extra pandoc
```

Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (6.5.4)
Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.9.4)
Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.12.3)
Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from nbconvert) (6.1.0)
Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.7.1)
Requirement already satisfied: entrypoints>=0.2.2 in

Campaign dataset@DhanunjayaReddy

August 30, 2024

1 Campaign Dataset

```
[125]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Downloading and reading the shopping csv file

```
[126]: df = pd.read_csv("campaign.csv")
df
```

```
[126]:
```

	ID	Year_Birth	Education	Marital_Status	Income	Kidhome	\
0	1826	1970	Graduation	Divorced	\$84,835.00	0	
1	1	1961	Graduation	Single	\$57,091.00	0	
2	10476	1958	Graduation	Married	\$67,267.00	0	
3	1386	1967	Graduation	Together	\$32,474.00	1	
4	5371	1989	Graduation	Single	\$21,474.00	1	
...	
2234	10142	1976	PhD	Divorced	\$66,476.00	0	
2235	5263	1977	2n Cycle	Married	\$31,056.00	1	
2236	22	1976	Graduation	Divorced	\$46,310.00	1	
2237	528	1978	Graduation	Married	\$65,819.00	0	
2238	4070	1969	PhD	Married	\$94,871.00	0	

	Teenhome	Dt_Customer	Recency	MntWines	...	NumCatalogPurchases	\
0	0	6/16/14	0	189	...		4
1	0	6/15/14	0	464	...		3
2	1	5/13/14	0	134	...		2
3	1	5/11/14	0	10	...		0
4	0	4/8/14	0	6	...		1
...	
2234	1	3/7/13	99	372	...		2
2235	0	1/22/13	99	5	...		0
2236	0	12/3/12	99	185	...		1
2237	0	11/29/12	99	267	...		4
2238	2	9/1/12	99	169	...		5

	NumStorePurchases	NumWebVisitsMonth	AcceptedCmp3	AcceptedCmp4	\
0	6	1	0	0	
1	7	5	0	0	
2	5	2	0	0	
3	2	7	0	0	
4	2	7	1	0	
...	
2234	11	4	0	0	
2235	3	8	0	0	
2236	5	8	0	0	
2237	10	3	0	0	
2238	4	7	0	0	1

	AcceptedCmp5	AcceptedCmp1	AcceptedCmp2	Complain	Country
0	0	0	0	0	SP
1	0	0	1	0	CA
2	0	0	0	0	US
3	0	0	0	0	AUS
4	0	0	0	0	SP
...
2234	0	0	0	0	US
2235	0	0	0	0	SP
2236	0	0	0	0	SP
2237	0	0	0	0	IND
2238	1	0	0	0	CA

[2239 rows x 27 columns]

```
[88]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2239 entries, 0 to 2238
Data columns (total 27 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    2239 non-null  int64
1   Year_Birth            2239 non-null  int64
2   Education              2239 non-null  object
3   Marital_Status        2239 non-null  object
4   Income                 2239 non-null  object
5   Kidhome                2239 non-null  int64
6   Teenhome               2239 non-null  int64
7   Dt_Customer            2239 non-null  object
8   Recency                2239 non-null  int64
9   MntWines               2239 non-null  int64
10  MntFruits               2239 non-null  int64
11  MntMeatProducts        2239 non-null  int64
```

12	MntFishProducts	2239	non-null	int64
13	MntSweetProducts	2239	non-null	int64
14	MntGoldProds	2239	non-null	int64
15	NumDealsPurchases	2239	non-null	int64
16	NumWebPurchases	2239	non-null	int64
17	NumCatalogPurchases	2239	non-null	int64
18	NumStorePurchases	2239	non-null	int64
19	NumWebVisitsMonth	2239	non-null	int64
20	AcceptedCmp3	2239	non-null	int64
21	AcceptedCmp4	2239	non-null	int64
22	AcceptedCmp5	2239	non-null	int64
23	AcceptedCmp1	2239	non-null	int64
24	AcceptedCmp2	2239	non-null	int64
25	Complain	2239	non-null	int64
26	Country	2239	non-null	object

dtypes: int64(22), object(5)

memory usage: 472.4+ KB

Unique number of values for specific categorical columns

```
[128]: columns_list = df[['ID', 'Education', 'Marital_Status', 'Country']]

for columns in columns_list.columns:
    unique_count = columns_list[columns].nunique()
    print(columns, "-", unique_count)
```

ID - 2239

Education - 5

Marital_Status - 8

Country - 8

Checking for the presence of null values in dataset.

```
[129]: df.isna().sum()
```

```
[129]: ID                0
Year_Birth             0
Education              0
Marital_Status         0
Income                0
Kidhome               0
Teenhome              0
Dt_Customer            0
Recency               0
MntWines              0
MntFruits             0
MntMeatProducts       0
MntFishProducts       0
MntSweetProducts      0
```

```

MntGoldProds      0
NumDealsPurchases 0
NumWebPurchases    0
NumCatalogPurchases 0
NumStorePurchases  0
NumWebVisitsMonth  0
AcceptedCmp3       0
AcceptedCmp4       0
AcceptedCmp5       0
AcceptedCmp1       0
AcceptedCmp2       0
Complain           0
Country            0
dtype: int64

```

shape of the dataset

```
[130]: df.shape
```

```
[130]: (2239, 27)
```

summary statistics of the dataset

```
[131]: df['Income'] = df['Income'].replace({'\': ' ', ': '}, regex=True).
        ↪astype(float)
df['Income'] = df['Income'].fillna(0).astype(int)
```

```
[132]: selected_variables = df[['Income', 'Kidhome',
                                'Teenhome', 'Recency', 'MntWines', 'MntFruits',
                                'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts',
                                'MntGoldProds', 'NumDealsPurchases', 'NumWebPurchases',
                                'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth']]

summary_df = selected_variables.describe()
summary_df
```

```
[132]:
```

	Income	Kidhome	Teenhome	Recency	MntWines	\
count	2239.000000	2239.000000	2239.000000	2239.000000	2239.000000	
mean	51412.792765	0.443948	0.506476	49.121036	304.067441	
std	22069.582225	0.538390	0.544555	28.963662	336.614830	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	34716.000000	0.000000	0.000000	24.000000	24.000000	
50%	51039.000000	0.000000	0.000000	49.000000	174.000000	
75%	68277.500000	1.000000	1.000000	74.000000	504.500000	
max	162397.000000	2.000000	2.000000	99.000000	1493.000000	

	MntFruits	MntMeatProducts	MntFishProducts	MntSweetProducts	\
count	2239.000000	2239.000000	2239.000000	2239.000000	


mean	26.307727	167.016525	37.538633	27.074587
std	39.781468	225.743829	54.637617	41.286043
min	0.000000	0.000000	0.000000	0.000000
25%	1.000000	16.000000	3.000000	1.000000
50%	8.000000	67.000000	12.000000	8.000000
75%	33.000000	232.000000	50.000000	33.000000
max	199.000000	1725.000000	259.000000	263.000000

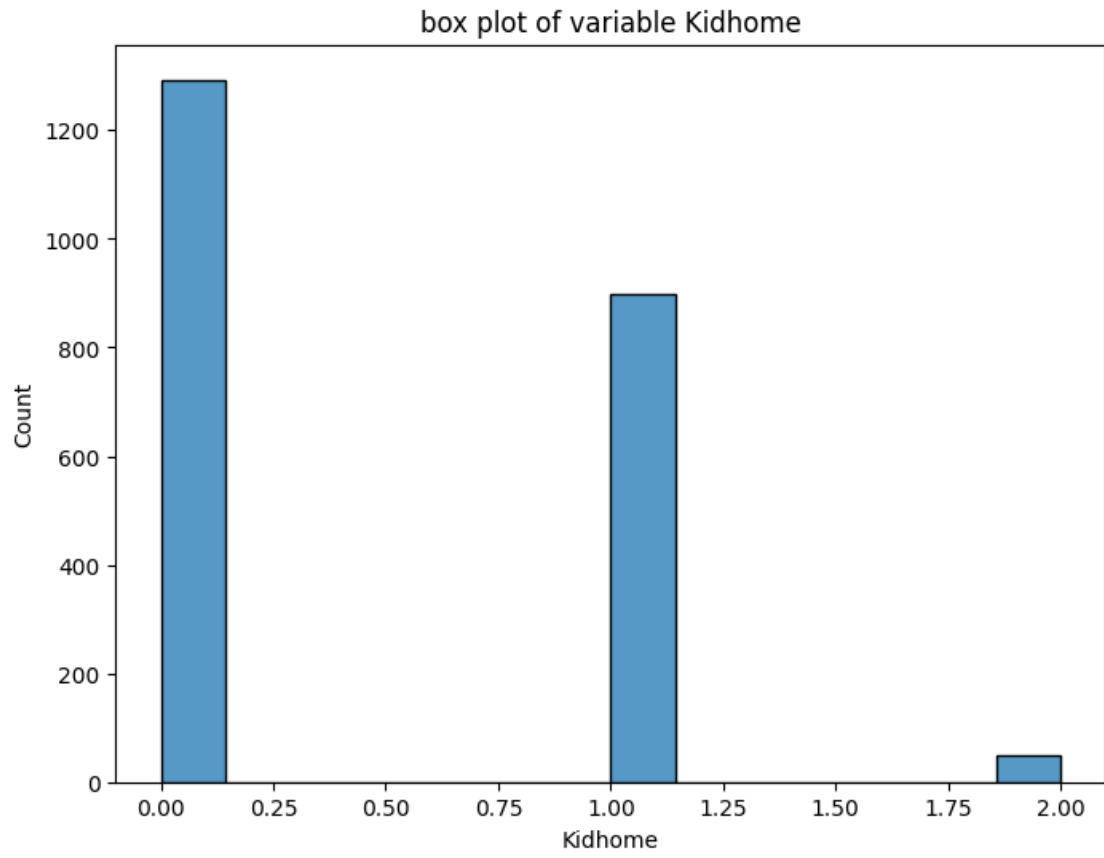
	MntGoldProds	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	\
count	2239.000000	2239.000000	2239.000000	2239.000000	
mean	44.036177	2.324252	4.085306	2.662796	
std	52.174700	1.932345	2.779240	2.923542	
min	0.000000	0.000000	0.000000	0.000000	
25%	9.000000	1.000000	2.000000	0.000000	
50%	24.000000	2.000000	4.000000	2.000000	
75%	56.000000	3.000000	6.000000	4.000000	
max	362.000000	15.000000	27.000000	28.000000	

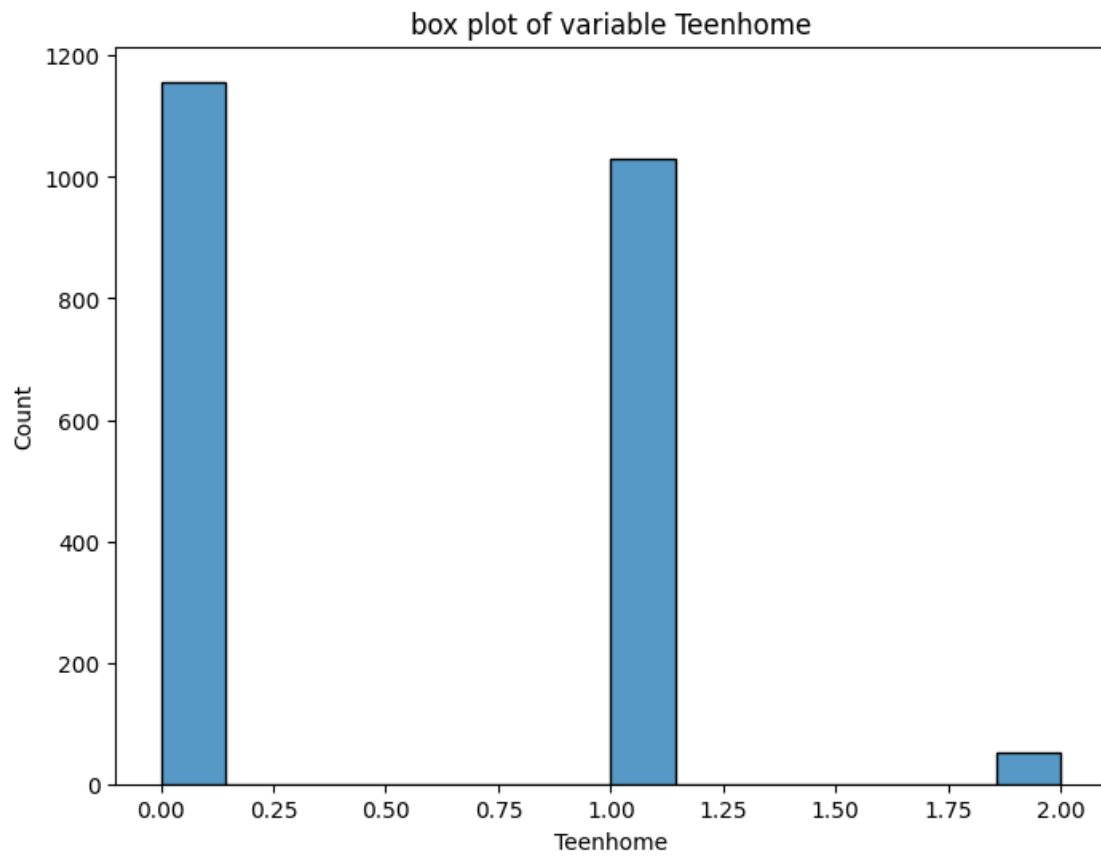
	NumStorePurchases	NumWebVisitsMonth
count	2239.000000	2239.000000
mean	5.791425	5.316213
std	3.251149	2.427144
min	0.000000	0.000000
25%	3.000000	3.000000
50%	5.000000	6.000000
75%	8.000000	7.000000
max	13.000000	20.000000

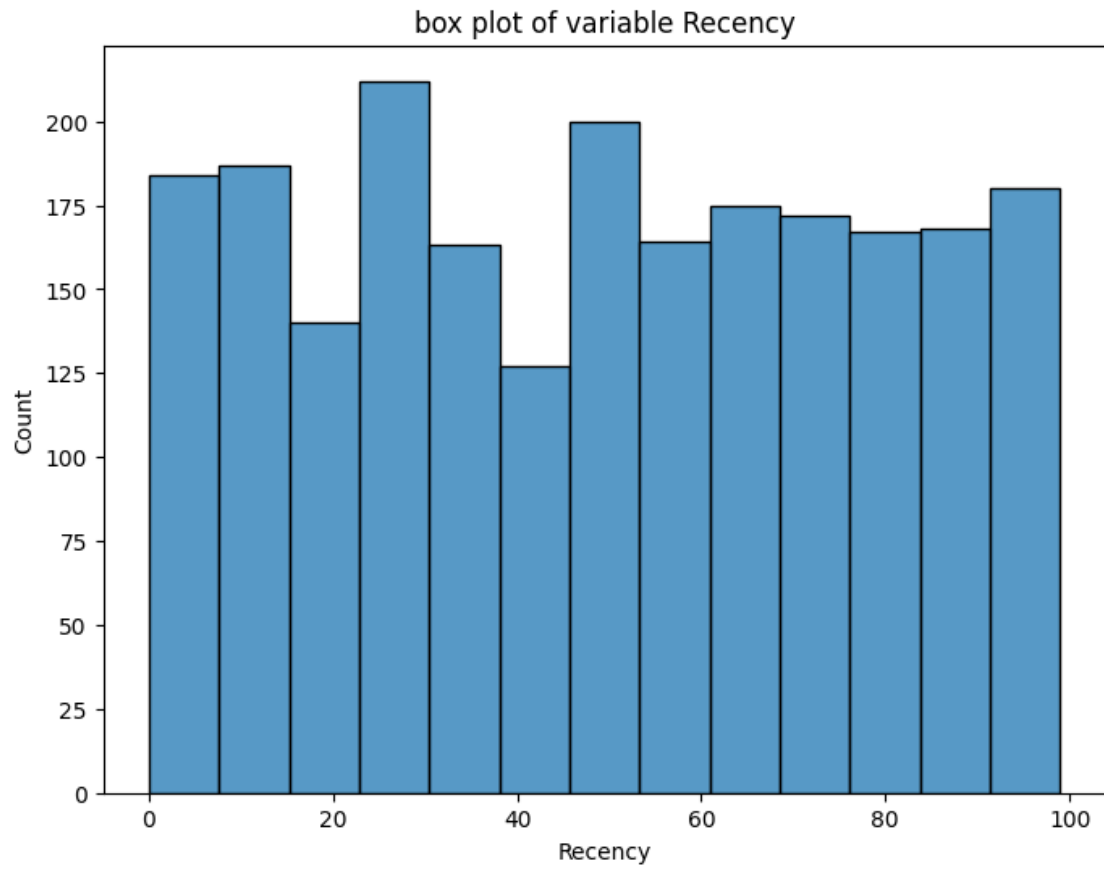
Distribution of the numerical features in the dataset

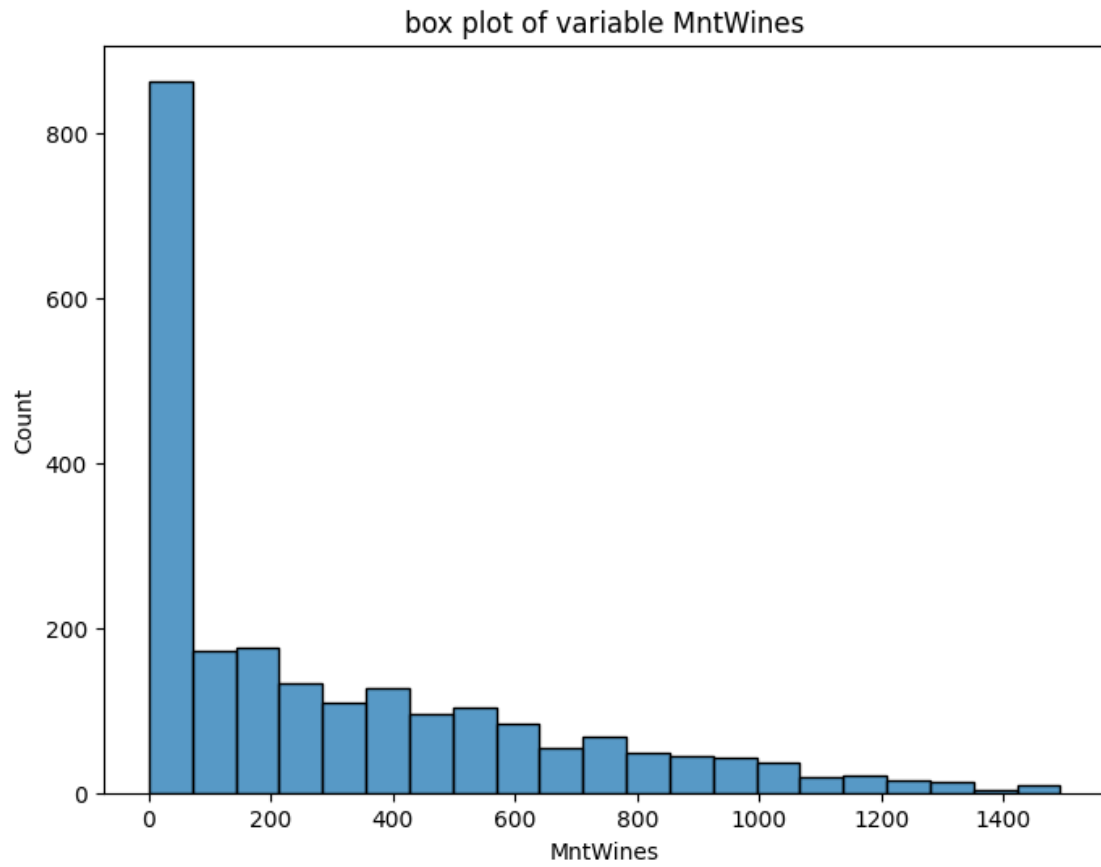
```
[133]: for variable in selected_variables:
        plt.figure(figsize = (8, 6))
        sns.histplot(data = selected_variables[variable])
        plt.title(f"box plot of variable {variable}")
```

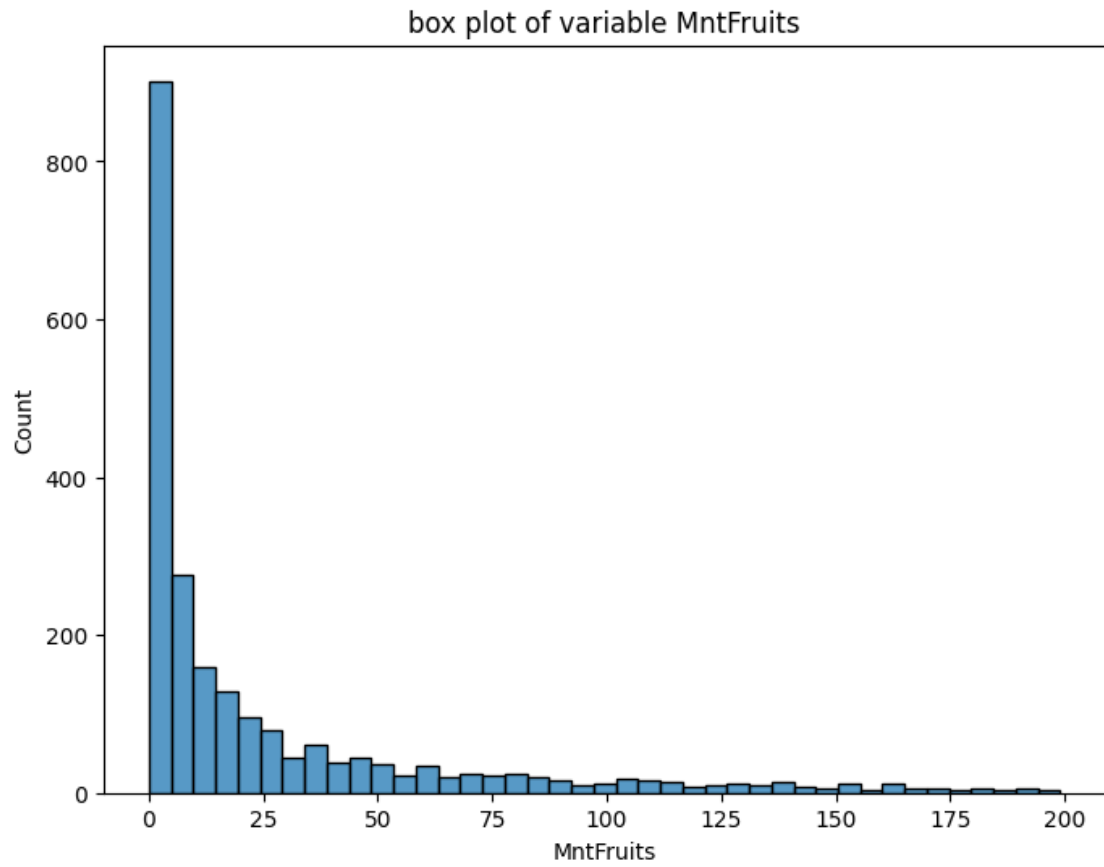


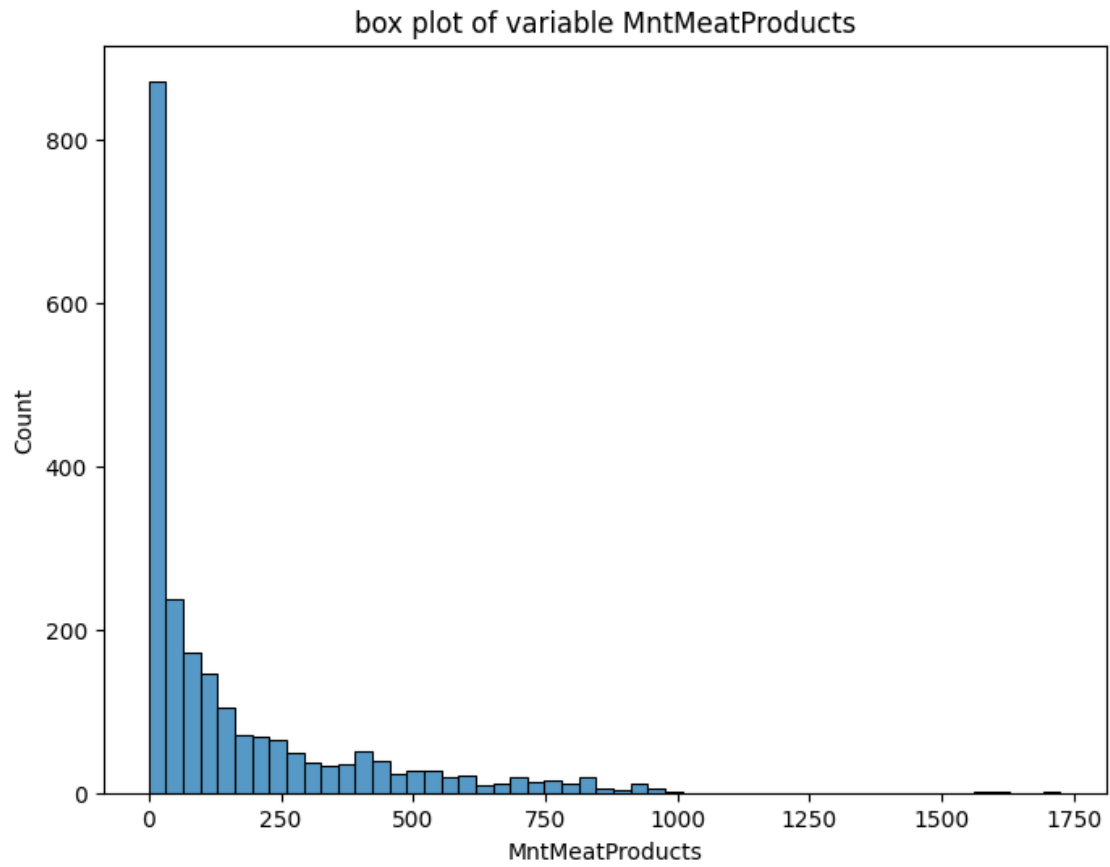


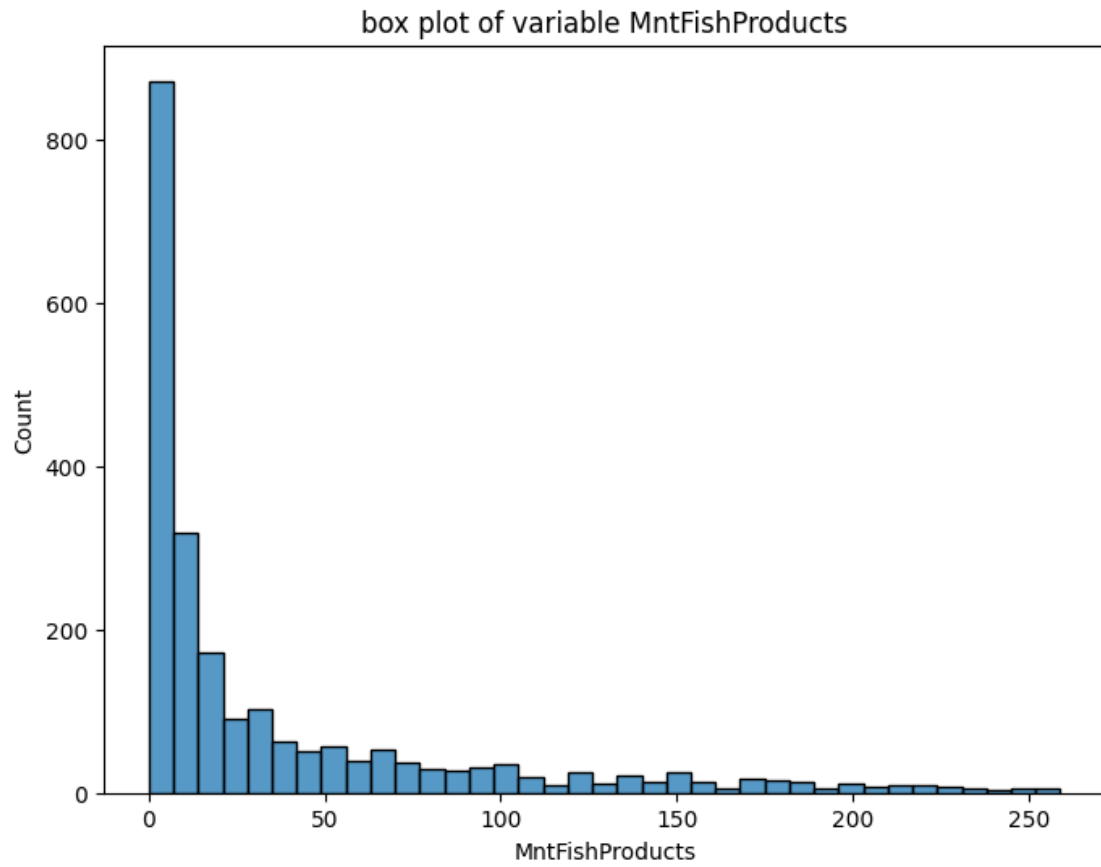











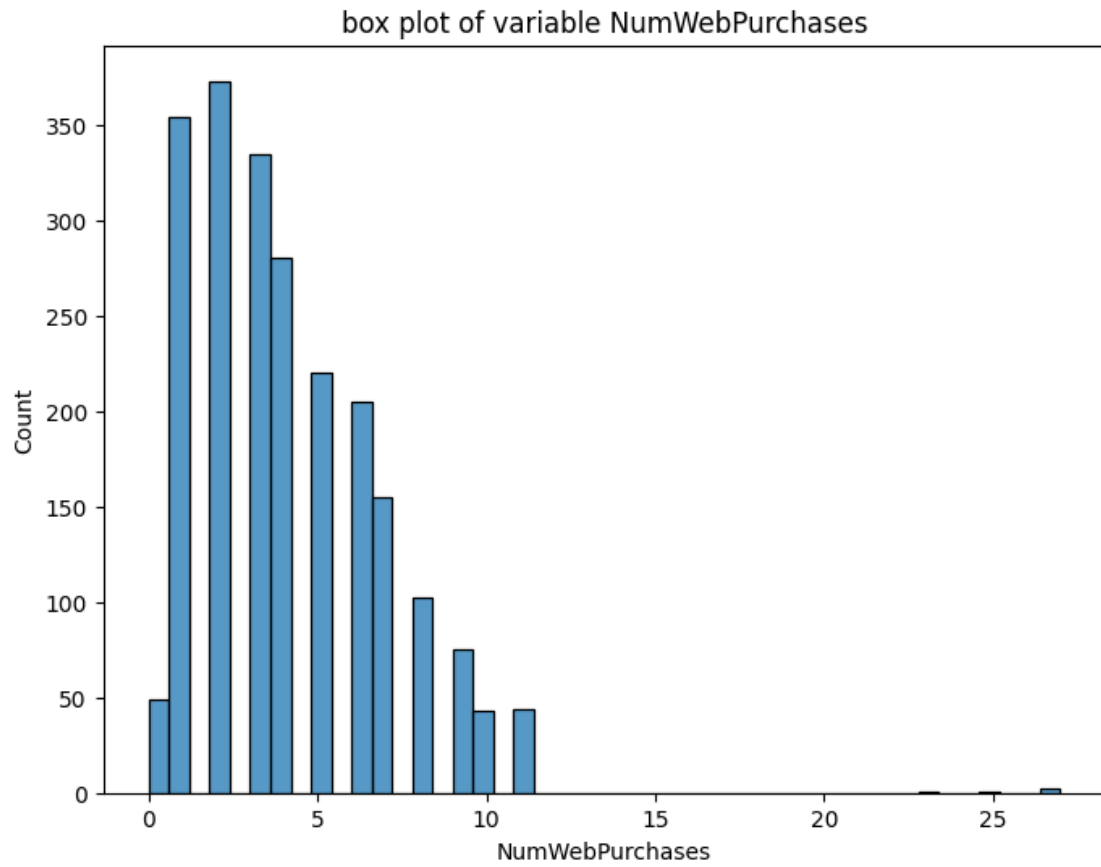


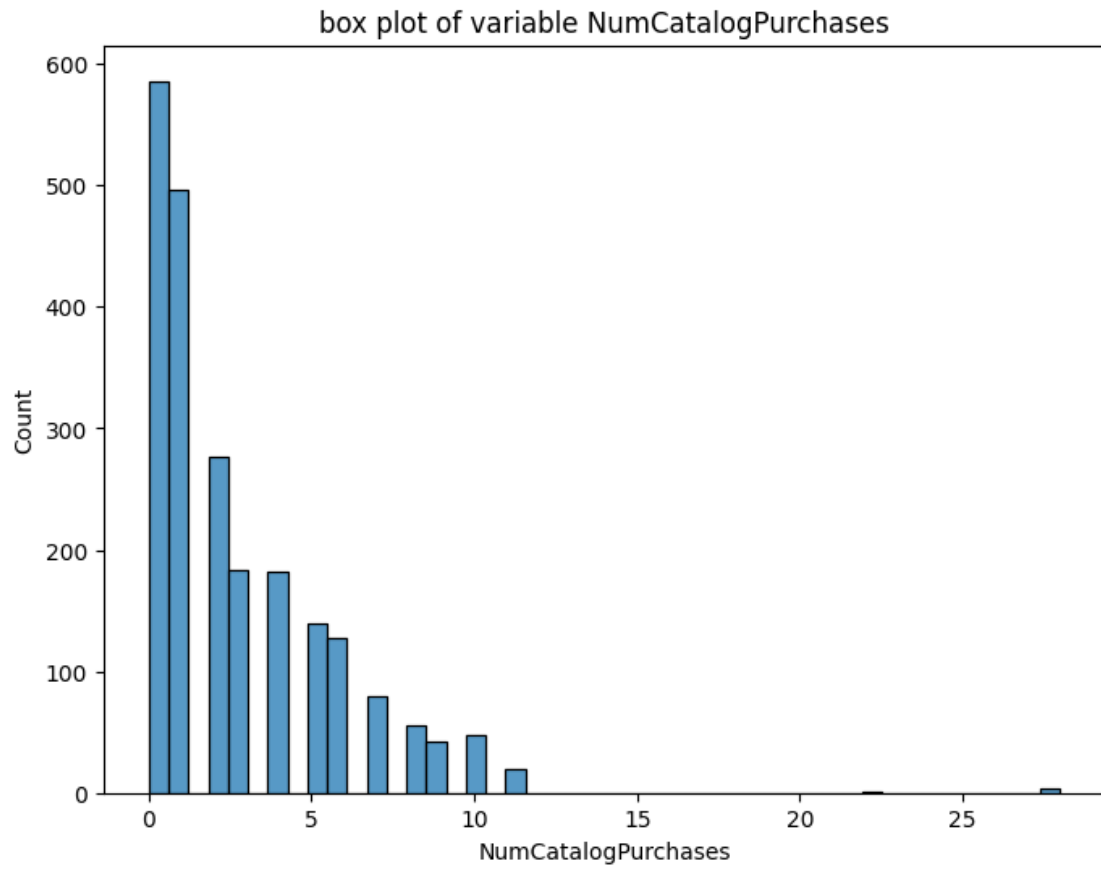


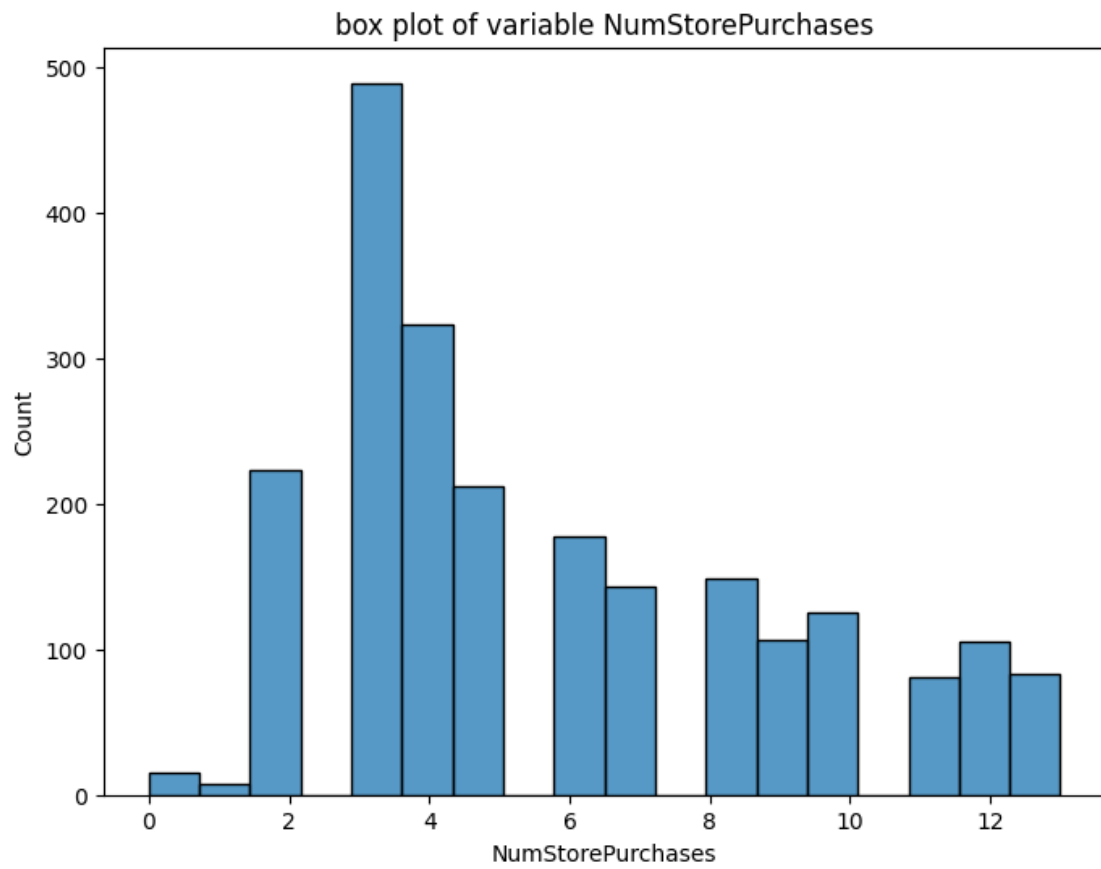




















Checking for the presence of outliers in the dataset


```
[97]: for variable in selected_variables:
      plt.figure(figsize = (8, 6))
      sns.boxplot(data = selected_variables[variable])
      plt.title(f"box plot of variable {variable}")
```





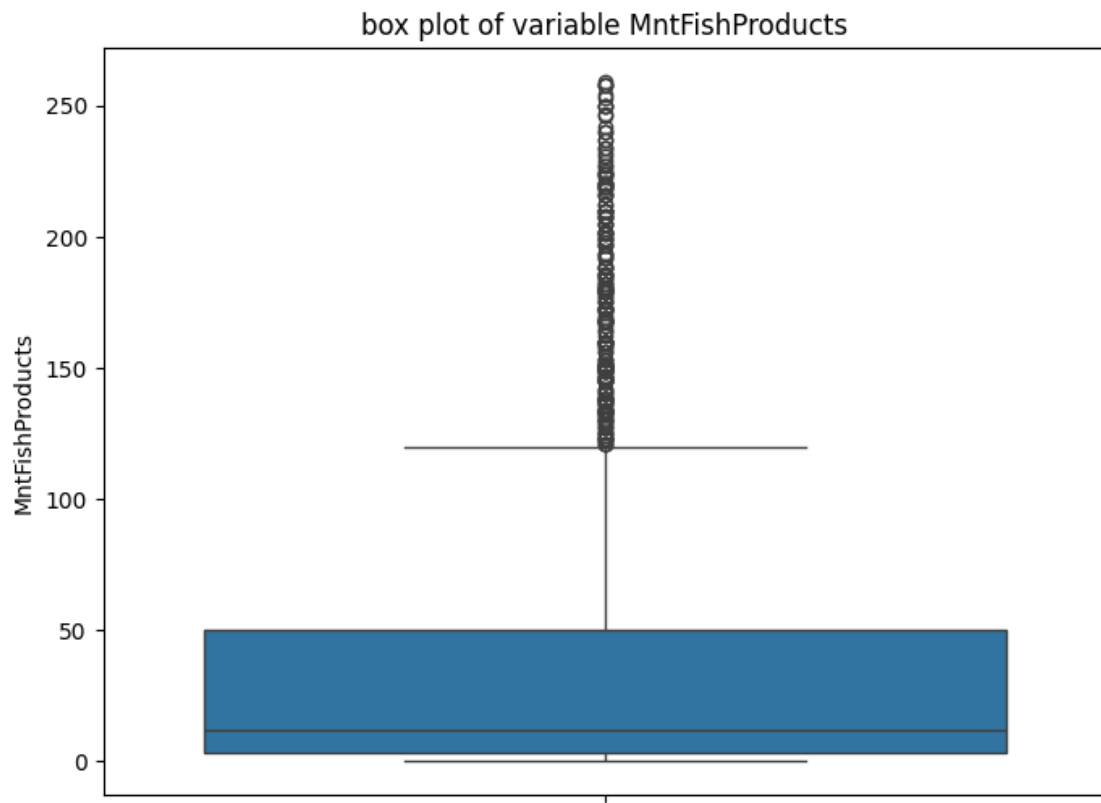














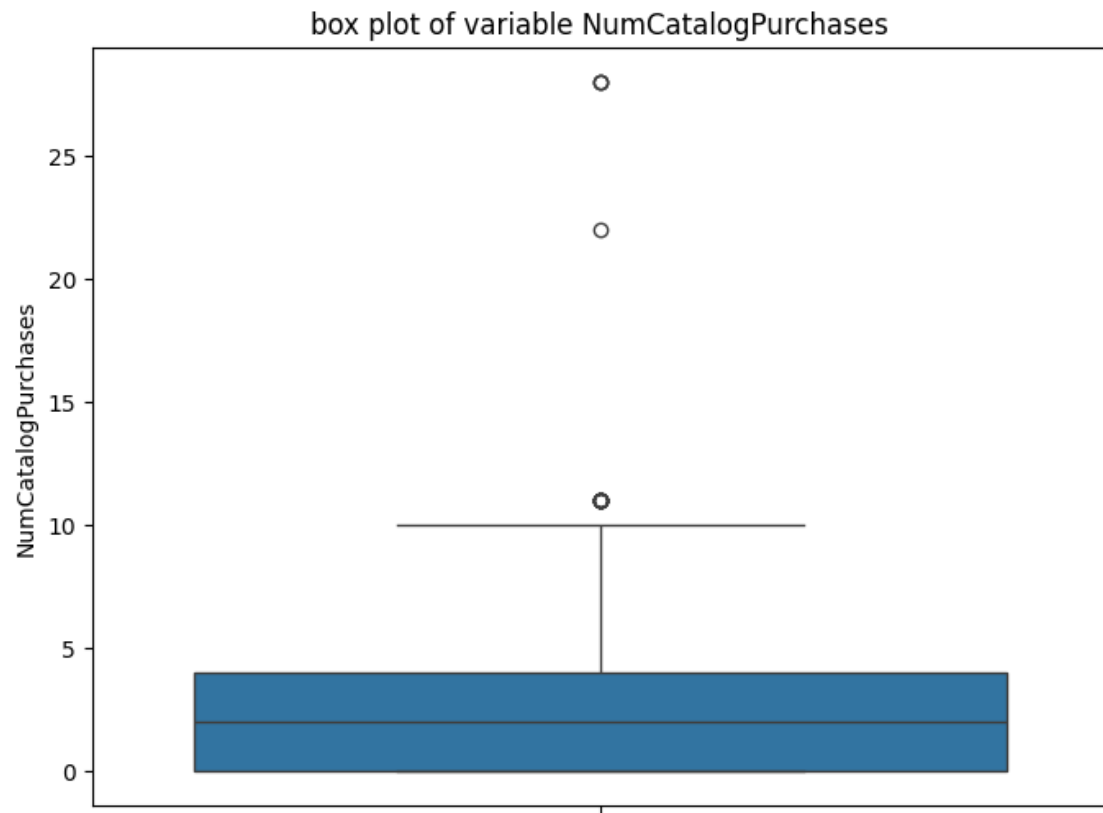





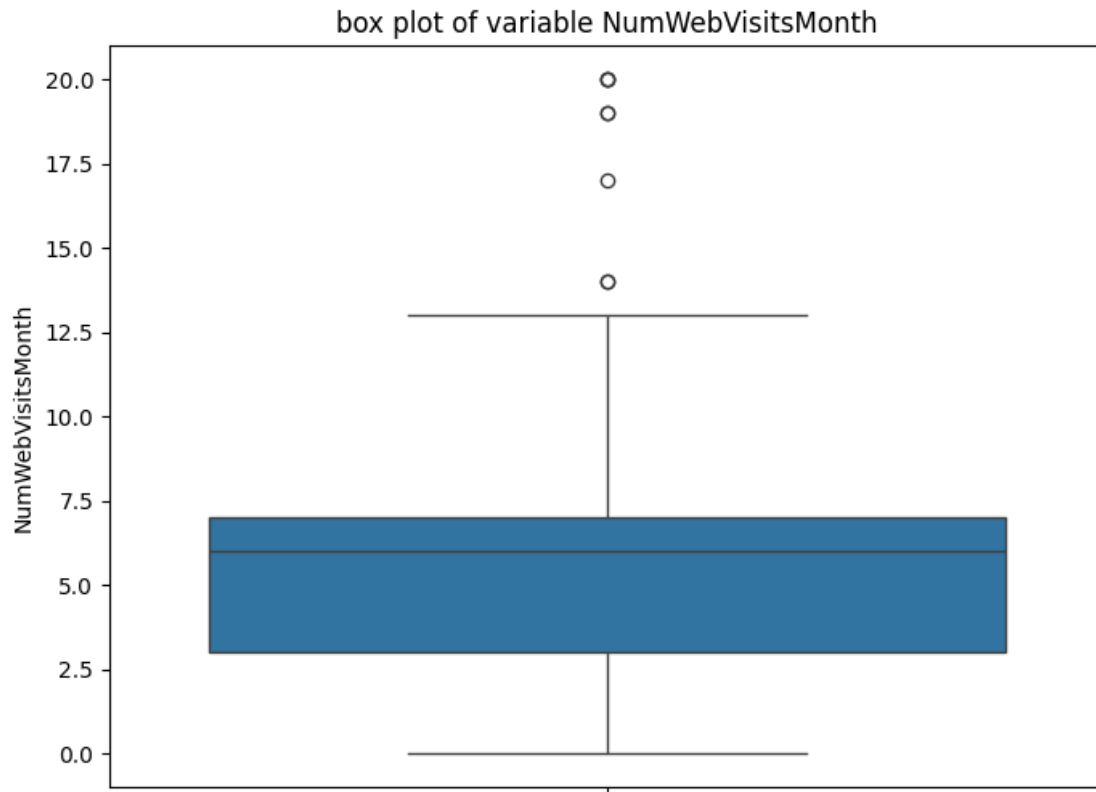












Calculating Total number of outliers

```
[134]: Q1 = summary_df.loc["25%"]
      Q3 = summary_df.loc["75%"]

      IQR = Q3- Q1
      print(IQR)
```

Income	33561.5
Kidhome	1.0
Teenhome	1.0
Recency	50.0
MntWines	480.5
MntFruits	32.0
MntMeatProducts	216.0
MntFishProducts	47.0
MntSweetProducts	32.0
MntGoldProds	47.0
NumDealsPurchases	2.0
NumWebPurchases	4.0
NumCatalogPurchases	4.0
NumStorePurchases	5.0

```
NumWebVisitsMonth          4.0
dtype: float64
```

```
[135]: lower_bound = Q1- 1.5*IQR
upper_bound = Q3 + 1.5*IQR

bounds_df = pd.DataFrame({"LowerBound" : lower_bound, "UpperBound":
    ↪upper_bound})
print(bounds_df)
```

	LowerBound	UpperBound
Income	-15626.25	118619.75
Kidhome	-1.50	2.50
Teenhome	-1.50	2.50
Recency	-51.00	149.00
MntWines	-696.75	1225.25
MntFruits	-47.00	81.00
MntMeatProducts	-308.00	556.00
MntFishProducts	-67.50	120.50
MntSweetProducts	-47.00	81.00
MntGoldProds	-61.50	126.50
NumDealsPurchases	-2.00	6.00
NumWebPurchases	-4.00	12.00
NumCatalogPurchases	-6.00	10.00
NumStorePurchases	-4.50	15.50
NumWebVisitsMonth	-3.00	13.00

```
[136]: outliers_lower = (summary_df < lower_bound).sum()
outliers_upper = (summary_df > upper_bound).sum()
total_outliers = outliers_lower + outliers_upper

ouliers_count_df = pd.DataFrame({"LowerBound_outliers" :outliers_lower,
    ↪"UpperBound_outliers" :outliers_upper, "Total" : total_outliers})
print(ouliers_count_df)
```

	LowerBound_outliers	UpperBound_outliers	Total
Income	0	1	1
Kidhome	0	1	1
Teenhome	0	1	1
Recency	0	1	1
MntWines	0	2	2
MntFruits	0	2	2
MntMeatProducts	0	2	2
MntFishProducts	0	2	2
MntSweetProducts	0	2	2
MntGoldProds	0	2	2
NumDealsPurchases	0	2	2
NumWebPurchases	0	2	2

NumCatalogPurchases	0	2	2
NumStorePurchases	0	1	1
NumWebVisitsMonth	0	2	2

Feature engineering

```
[137]: bins = [0, 5000, 25000, 45000, 65000, 85000, 105000, 125000, 145000, 165000]

labels = ['<=5k', '>5k-25k', '>25k-45k', '>45k-65k', '>65k-85k', '>85k-105k',
↪ '>105k-125k', '>125k-145k', '>145-165k']

df['Income_labels'] = pd.cut(df['Income'], bins = bins, labels = labels)
```

```
[138]: import datetime
df['Age'] = datetime.datetime.now().year - df['Year_Birth']
```

```
[139]: bins = [25, 45, 65, 85, 105, 125, 135]

labels = ['25-45', '>45-65', '>65-85', '>85-105', '>105-125', '>125+']

df['age_labels'] = pd.cut(df['Age'], bins = bins, labels = labels)
```

```
[140]: df['Dt_Customer'] = pd.to_datetime(df['Dt_Customer'])
df['Customer_period'] = (datetime.datetime.now().year - df['Dt_Customer'].dt.
↪ year) * 12 + (datetime.datetime.now().month - df['Dt_Customer'].dt.month)
```

```
[141]: df['Totalamt_spent'] = df['MntWines'] + df['MntFruits'] + df['MntMeatProducts'] +
↪ df['MntFishProducts'] + df['MntSweetProducts'] + df['MntGoldProds']

df['Total_purchases'] = df['NumDealsPurchases'] + df['NumWebPurchases'] +
↪ df['NumCatalogPurchases'] + df['NumStorePurchases']
```

Total amt spent on different products categorized under income labels.

```
[142]: import warnings

warnings.filterwarnings('ignore', category=FutureWarning)
warnings.filterwarnings('ignore')

Income_label_amt_spent = df.groupby('Income_labels')[['MntWines', 'MntFruits',
↪ 'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
↪ 'Totalamt_spent']].sum()
Income_label_amt_spent = Income_label_amt_spent.sort_values(by =
↪ ['Totalamt_spent'], ascending = False).reset_index()
print(Income_label_amt_spent)


sns.barplot(data = Income_label_amt_spent, x = 'Income_labels', y =
↪ 'Totalamt_spent', color = 'y')
```



```
plt.xticks(rotation = 90)
plt.show()
```

	Income_lables	MntWines	MntFruits	MntMeatProducts	MntFishProducts	\
0	>65k-85k	350199	32641	215341	47866	
1	>45k-65k	215973	14038	73676	18969	
2	>85k-105k	66082	5841	49390	7832	
3	>25k-45k	39914	4191	21323	6591	
4	>5k-25k	2661	1467	3508	1904	
5	>145-165k	203	22	4957	26	
6	<=5k	27	8	1743	6	
7	>105k-125k	1015	183	107	203	
8	>125k-145k	0	0	0	0	

	MntSweetProducts	MntGoldProds	Totalamt_spent
0	33792	41565	721404
1	13547	32050	368253
2	6624	5705	141474
3	4109	13304	89432
4	1524	4237	15301
5	9	18	5235
6	7	326	2117
7	283	210	2001
8	0	0	0



Type of purchases categorized under income labels.


```
[143]: Income_label_purchases = df.groupby('Income_labels')[['NumDealsPurchases',
    ↳ 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases',
    ↳ 'Total_purchases']].sum()
Income_label_purchases = Income_label_purchases.sort_values(by =
    ↳ ['Total_purchases'], ascending = False).reset_index()
print(Income_label_purchases)

sns.barplot(data = Income_label_purchases, x = 'Income_labels', y =
    ↳ 'Total_purchases', color = 'y')
plt.xticks(rotation = 90)
plt.show()
```

	Income_labels	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	\
0	>65k-85k	970	3122	3072	
1	>45k-65k	2024	3271	1646	
2	>25k-45k	1510	1660	439	

3	>85k-105k	75	467	548
4	>5k-25k	491	468	99
5	>145-165k	30	1	78
6	<=5k	45	25	28
7	>105k-125k	0	36	8
8	>125k-145k	0	0	0

	NumStorePurchases	Total_purchases
0	4945	12109
1	4229	11170
2	2292	5901
3	708	1798
4	662	1720
5	3	112
6	0	98
7	13	57
8	0	0



Purchases categorized for number of teenagers and kids in each household

```
[144]: df.groupby('Teenhome')[['Totalamt_spent', 'Total_purchases']].sum().
        ↪reset_index()
```

```
[144]:   Teenhome  Totalamt_spent  Total_purchases
0         0         802199         16061
1         1         524091         16338
2         2          30636           881
```

```
[145]: df.groupby('Kidhome')[['Totalamt_spent', 'Total_purchases']].sum().reset_index()
```

```
[145]:   Kidhome  Totalamt_spent  Total_purchases
0         0        1165330         23395
1         1        184624          9416
2         2          6972           469
```

Total amt spent on different products categorized under Age labels.

```
[146]: import warnings

warnings.filterwarnings('ignore', category=FutureWarning)
warnings.filterwarnings('ignore')


Age_label_amt_spent = df.groupby('age_labels')[['MntWines', 'MntFruits',
        ↪'MntMeatProducts', 'MntFishProducts', 'MntSweetProducts', 'MntGoldProds',
        ↪'Totalamt_spent']].sum()
Age_label_amt_spent = Age_label_amt_spent.sort_values(by = ['Totalamt_spent'],
        ↪ascending = False).reset_index()
print(Age_label_amt_spent)

sns.barplot(data = Age_label_amt_spent, x = 'age_labels', y = 'Totalamt_spent',
        ↪color = 'b')
plt.xticks(rotation = 90)
plt.show()
```

	age_labels	MntWines	MntFruits	MntMeatProducts	MntFishProducts	\
0	>45-65	367850	30701	187090	42260	
1	>65-85	196310	14039	100228	22594	
2	25-45	115869	14013	86057	19077	
3	>105-125	770	150	570	111	
4	>125+	8	0	5	7	
5	>85-105	0	0	0	0	

	MntSweetProducts	MntGoldProds	Totalamt_spent
0	31383	51898	711182
1	15432	26147	374750
2	13737	20301	269054

3	68	249	1918
4	0	2	22
5	0	0	0



Type of purchases categorized under Age labels.


```
[147]: Age_label_purchases = df.groupby('age_labels')[['NumDealsPurchases',
↳ 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases',
↳ 'Total_purchases']].sum()
Age_label_purchases = Age_label_purchases.sort_values(by = ['Total_purchases'],
↳ ascending = False).reset_index()
print(Age_label_purchases)

sns.barplot(data = Age_label_purchases, x = 'age_labels', y =
↳ 'Total_purchases', color = 'b')
plt.xticks(rotation = 90)
plt.show()
```

age_labels	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	\
------------	-------------------	-----------------	---------------------	---

0	>45-65	3097	5068	3101
1	>65-85	1213	2395	1713
2	25-45	891	1677	1141
3	>105-125	2	6	7
4	>125+	1	1	0
5	>85-105	0	0	0

	NumStorePurchases	Total_purchases
0	7089	18355
1	3331	8652
2	2539	6248
3	6	21
4	2	4
5	0	0



Type of purchases categorized under different Education levels.

[148]:

```

Purchases_by_Education_level = df.groupby('Education')[['NumDealsPurchases',
↳ 'NumWebPurchases', 'NumCatalogPurchases', 'NumStorePurchases',
↳ 'NumWebVisitsMonth']].sum()
Purchases_by_Education_level = Purchases_by_Education_level.reset_index()
print(Purchases_by_Education_level)


purchase_types = ['NumDealsPurchases', 'NumWebPurchases',
↳ 'NumCatalogPurchases', 'NumStorePurchases', 'NumWebVisitsMonth']


for purchase_type in purchase_types:
    sns.barplot(data = Purchases_by_Education_level, x = 'Education', y =
↳ purchase_type, color = 'grey')
    plt.title(purchase_type)
    plt.xlabel('Education level')
    plt.show()


```


	Education	NumDealsPurchases	NumWebPurchases	NumCatalogPurchases	\
0	2n Cycle	456	757	471	
1	Basic	97	102	26	
2	Graduation	2599	4646	3071	
3	Master	898	1492	951	
4	PhD	1154	2150	1443	


	NumStorePurchases	NumWebVisitsMonth
0	1118	1107
1	154	371
2	6567	5953
3	2182	1916
4	2946	2556











Number of customers attracted for each Different campaign

```
[149]: count_of_offers = df[['AcceptedCmp1', 'AcceptedCmp2', 'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5']].sum()
count_of_offers = count_of_offers.reset_index()
count_of_offers.columns = ['Campaign', 'count']
count_of_offers
```

```
[149]:
```

	Campaign	count
0	AcceptedCmp1	144
1	AcceptedCmp2	30
2	AcceptedCmp3	163
3	AcceptedCmp4	167
4	AcceptedCmp5	163

Total number of Complaints


```
[150]: Total_complaints = df['Complain'].sum()
Total_complaints
```

[150]: 21

Total amt spent with respect to customer period. (in months)

```
[151]: customer_purchases_wtr_period = df.  
        ↳groupby('Customer_period')[['Totalamt_spent', 'Total_purchases']].sum()  
customer_purchases_wtr_period = customer_purchases_wtr_period.sort_values(by =  
        ↳['Totalamt_spent'], ascending = False).reset_index()  
print(customer_purchases_wtr_period)  
  
sns.lineplot(data=customer_purchases_wtr_period, x='Customer_period',  
        ↳y='Totalamt_spent', label='Total Amount Spent')  
plt.title('Customer Amount Spent Over Time')  
plt.xlabel('Customer Period (in months)')  
plt.show()
```

	Customer_period	Totalamt_spent	Total_purchases
0	144	92046	2026
1	143	85492	1812
2	141	77190	1722
3	135	72954	1593
4	139	71351	1789
5	138	68879	1671
6	137	67322	1667
7	130	66221	1704
8	142	64802	1598
9	129	60064	1394
10	136	58135	1369
11	132	55393	1478
12	123	54648	1475
13	133	49986	1358
14	128	49455	1243
15	127	48932	1306
16	134	48171	1335
17	131	47432	1236
18	140	47382	1116
19	125	47247	1306
20	124	44938	1154
21	126	42970	972
22	122	34559	911
23	145	1357	45




Total amt spent and mean amt spent for each country

```
[152]: Country_purchases = df.groupby('Country')['Totalamt_spent'].aggregate({'sum',
    ↳ 'mean'})
Country_purchases = Country_purchases.sort_values(by = ['mean'], ascending =
    ↳ False).reset_index()
Country_purchases.columns = ['country', 'Total_amt_spent', 'mean_amt_spent']
print(Country_purchases)

Country_purchases.plot(kind='barh', x='country', y='Total_amt_spent',
    ↳ figsize=(12, 8), color = 'r')
plt.title('Total and Mean Amount Spent by Country')
plt.xlabel('Amount Spent')
plt.ylabel('Country')
plt.show()
```

	country	Total_amt_spent	mean_amt_spent
0	ME	3122	1040.666667
1	CA	168532	628.850746
2	SA	211009	628.002976

3	GER	74913	624.275000
4	US	67882	622.770642
5	SP	662220	604.767123
6	AUS	89763	561.018750
7	IND	79485	537.060811



Hypothesis testing

Is income of customers dependent on their education

```
[153]: from scipy.stats import f_oneway

alpha = 0.05
# H0: The means of Incomes between different education levels would be equal
# H1: The means of Incomes between different education levels would be not_
↳equal

education_groups = [group['Income'].values for name, group in df.
↳groupby('Education')]

f_statistic, p_value = f_oneway(*education_groups)
print(f_statistic, p_value)
print("<----->")
```

```

if p_value < 0.05:
    print("Reject the null hypothesis. Income depends on education level.")
else:
    print("Fail to reject null hypothesis : Income doesn't depend on Education_
↪levels.")

```

35.42763066856272 9.87796950058819e-29

<----->

Reject the null hypothesis. Income depends on education level.

Do higher income people spend more (take in account spending in all categories together)

```

[154]: import scipy.stats as stats

corr_coefficient, p_value = stats.pearsonr(df['Income'], df['Totalamt_spent'])

print(f"Pearson correlation coefficient: {corr_coefficient}")
print(f"P-value: {p_value}")

print("<----->")

if p_value < 0.05:
    print("Reject the null hypothesis. There is a significant linear_
↪relationship between income and purchases.")
else:
    print("Fail to reject the null hypothesis. No significant linear_
↪relationship exists between income and purchases.")

```

Pearson correlation coefficient: 0.7706290398754154

P-value: 0.0

<----->

Reject the null hypothesis. There is a significant linear relationship between income and purchases.

Do couples spend more or less money on wine than people living alone (set 'Married', 'Together': 'In couple' and 'Divorced', 'Single', 'Absurd', 'Widow', 'YOLO': 'Alone')

```

[155]: df['Living_Status'] = df['Marital_Status'].apply(lambda x: 'Couple' if x in_
↪['Married', 'Together'] else 'Alone')

wine_spending_couples = df[df['Living_Status'] == 'Couple']['MntWines']
wine_spending_alone = df[df['Living_Status'] == 'Alone']['MntWines']

import scipy.stats as stats

# Perform the t-test

```



```

t_statistic, p_value = stats.ttest_ind(wine_spending_couples,
    ↪ wine_spending_alone, equal_var=False) # Use equal_var=False if variances
    ↪ are unequal

print(f"T-statistic: {t_statistic}")
print(f"P-value: {p_value}")

print("<----->")

if p_value < 0.05:
    print("Reject the null hypothesis. There is a significant difference in
    ↪ wine spending between couples and people living alone.")
else:
    print("Fail to reject the null hypothesis. No significant difference in
    ↪ wine spending exists between couples and people living alone.")

```

T-statistic: -0.2711337908368919

P-value: 0.7863223090103292

<----->

Fail to reject the null hypothesis. No significant difference in wine spending exists between couples and people living alone.

Are people with lower income are more attracted towards campaign or simply put accept more campaigns. (create two income brackets one below median , other above median income and create a column which tells if they have ever accepted any campaign)

```

[156]: median_income = df['Income'].median()

df['Income_Bracket'] = df['Income'].apply(lambda x: 'Below_Median' if x <
    ↪ median_income else 'Above_Median')

df['Accepted_Any_Campaign'] = df[['AcceptedCmp1', 'AcceptedCmp2',
    ↪ 'AcceptedCmp3', 'AcceptedCmp4', 'AcceptedCmp5']].sum(axis=1).apply(lambda x:
    ↪ 1 if x > 0 else 0)

import scipy.stats as stats

# Create a contingency table
contingency_table = pd.crosstab(df['Income_Bracket'],
    ↪ df['Accepted_Any_Campaign'])

# Perform the Chi-Square test
chi2, p, dof, expected = stats.chi2_contingency(contingency_table)

print(f"Chi-Square statistic: {chi2}")
print(f"P-value: {p}")

```

```

print("<----->")

if p < 0.05:
    print("Reject the null hypothesis. There is a significant association_
    ↳between income level and campaign acceptance.")
else:
    print("Fail to reject the null hypothesis. No significant association_
    ↳between income level and campaign acceptance.")

```

Chi-Square statistic: 138.8199834041559

P-value: 4.8224046007539564e-32

<----->

Reject the null hypothesis. There is a significant association between income level and campaign acceptance.

```

[157]: from google.colab import drive
drive.mount("/content/drive")

```

Mounted at /content/drive

```

[ ]: !pip install nbconvert

!apt-get install texlive texlive-xetex texlive-latex-extra pandoc

!jupyter nbconvert --to pdf "/content/drive/MyDrive/Colab Notebooks/Campaign_
↳dataset@DhanunjayaReddy.ipynb".ipynb

```

Requirement already satisfied: nbconvert in /usr/local/lib/python3.10/dist-packages (6.5.4)

Requirement already satisfied: lxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.9.4)

Requirement already satisfied: beautifulsoup4 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (4.12.3)

Requirement already satisfied: bleach in /usr/local/lib/python3.10/dist-packages (from nbconvert) (6.1.0)

Requirement already satisfied: defusedxml in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.7.1)

Requirement already satisfied: entrypoints>=0.2.2 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.4)

Requirement already satisfied: Jinja2>=3.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (3.1.4)

Requirement already satisfied: jupyter-core>=4.7 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (5.7.2)

Requirement already satisfied: jupyterlab-pygments in /usr/local/lib/python3.10/dist-packages (from nbconvert) (0.3.0)

Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from nbconvert) (2.1.5)